

Artificial Intelligence

Algorithms and Applications with Python

Chapter 9



Dr. Dominik Jung

dominik.jung42@gmail.com



python

Outline

9 Language and Image Processing

9.1 Applied Natural Language Processing

9.2 Speech Processing and Communication

9.3 Perception and Image Processing

9.4 Generative Modelling

Lectorial 7: Building NLP and Machine Learning Pipelines for Webservices

► **What we will learn:**

- We will discuss how to make use of the copious knowledge that is expressed in natural language
- And how to process and prepare language data for machines to build state-of-the-art machine learning pipelines for language processing
- We learn new type of models, so called generative models, which allow to generate new instances based on the knowledge of your model



Image source: [↗ Pixabay](#) (2021) / [↗ CC0](#)

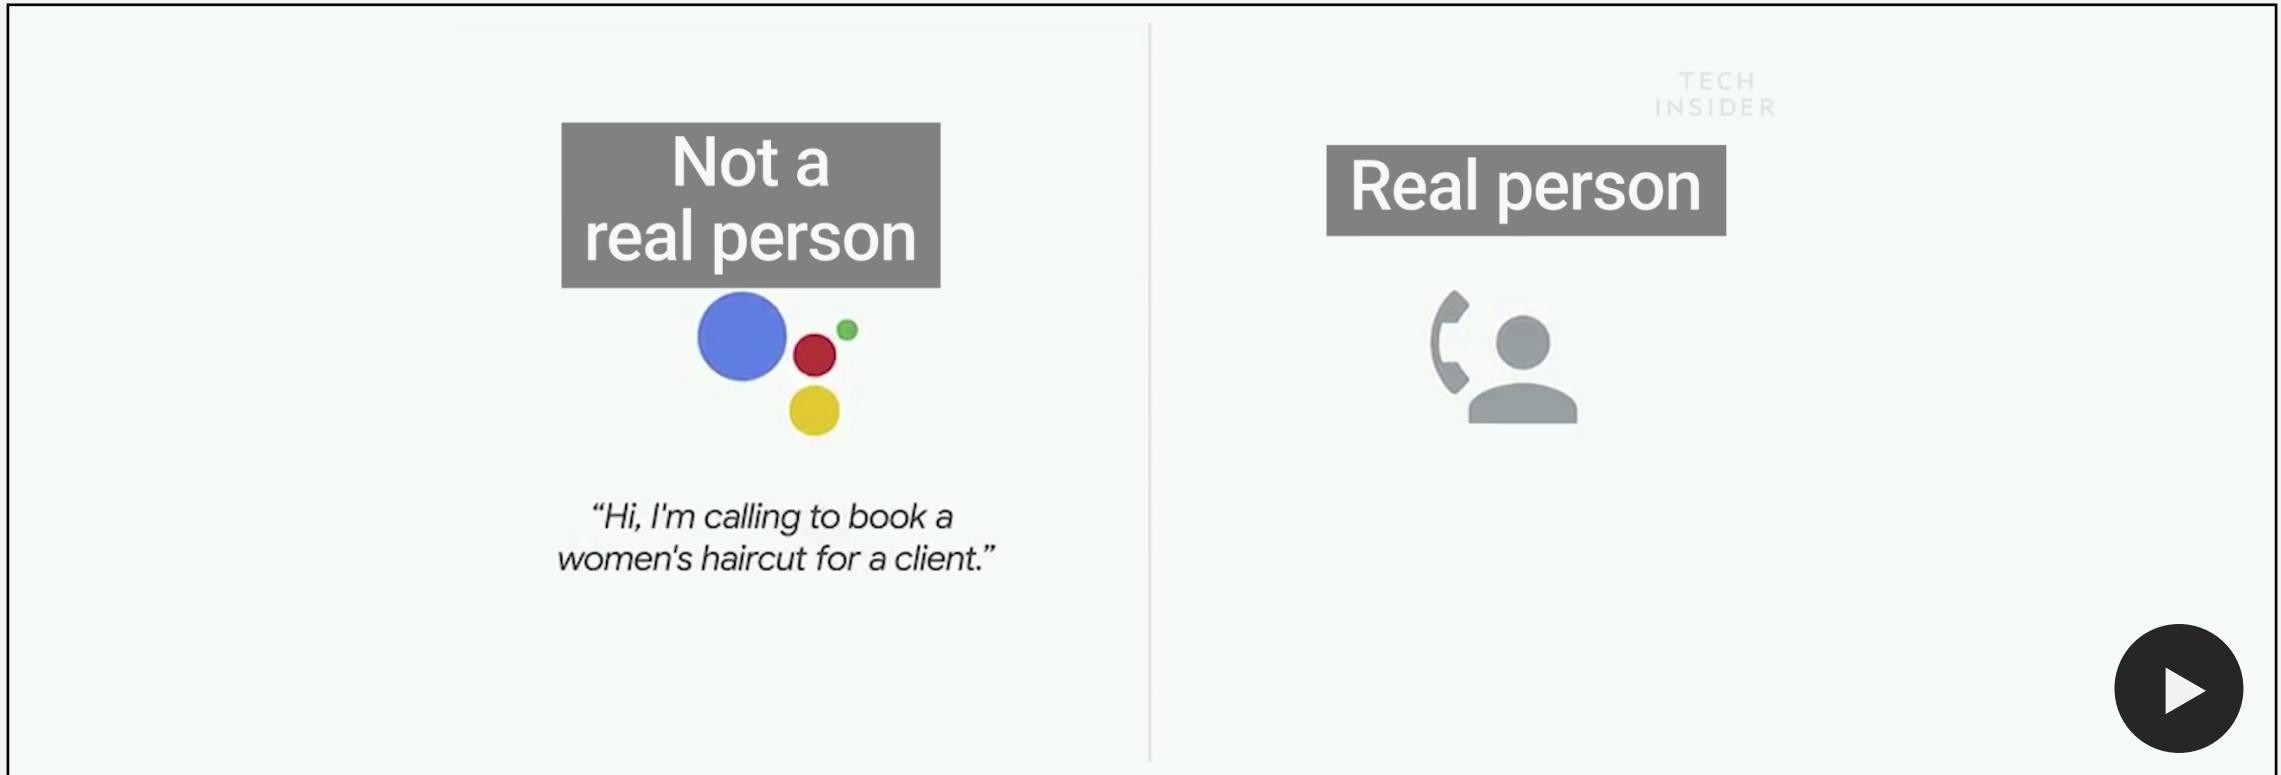
► **Duration:**

- 135 min + 90 (Lectorial)

► **Relevant for Exam:**

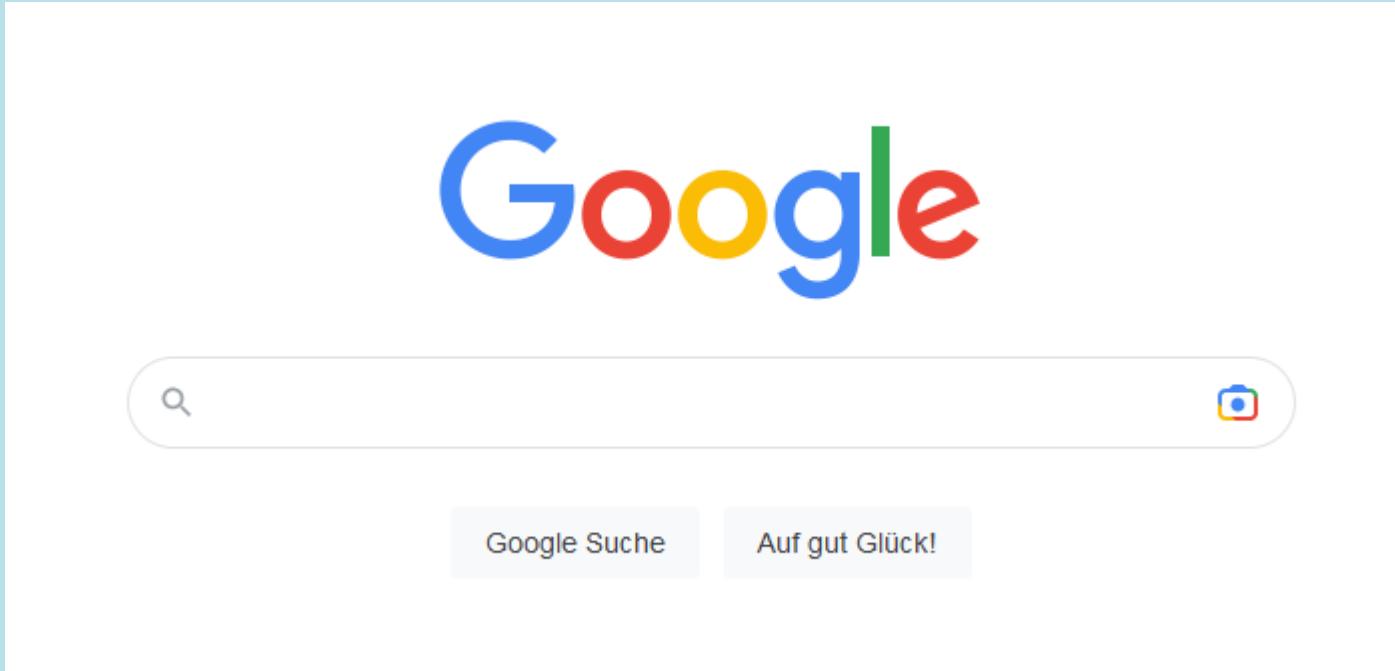
- 9.1, 9.4

9.1 Natural Language Processing



Natural Language Processing is the branch of computer science focused on developing systems that allow computers to communicate with people using everyday language

9.1 Application: Information Retrieval and Search Engines



D

Information Retrieval is the science building a system for recovering specific information from stored data base based on user's needs.

9.1 Text Processing Example I

The screenshot shows the Wikipedia page for 'Porsche'. The page title is 'Porsche'. The main content starts with a summary: 'This article is about the automotive brand and manufacturer, Porsche AG. For the holding company that is the majority owner of Volkswagen Group, see [Porsche SE](#). For other uses of Porsche, see [Porsche \(disambiguation\)](#)'. Below this, a section discusses the company's history and products. The company name 'Dr.-Ing. h.c. F. Porsche AG' is followed by the text 'usually shortened to Porsche (German pronunciation: [ˈpoʁʃə] (listen); see below)', which is highlighted with a red box. To the right of the text is the official Porsche logo, featuring the crest and the word 'PORSCHE'.

Preprocessing

- Extract plain text
- Reduce complexity
- prepare transformation

...
<p>Dr.-Ing. h.c. F. Porsche AG,
usually shortened to Porsche
<small>German
pronunciation: </small><span
title="Representation in the International
Phonetic Alphabet (IPA)" class="IPA"><a
href="/wiki/Help:IPA/Standard_German"
title="Help:IPA/Standard
German">['poʁʃə] <span
class="nowrap" style="font-size:85%">(<span
class="unicode haudio"><span
style="white-space:nowrap; margin-
right:.25em;"><a href="/wiki/File:De-
Porsche.ogg" title="About this sound">

9.1 Text Processing Example II

- Remove html tags

Dr.-Ing. h.c. F. Porsche AG, usually shortened to Porsche (German pronunciation: [ˈpoʃə] (audio speaker iconlisten); see below), is a German automobile manufacturer specializing in high-performance sports cars, SUVs and sedans, headquartered in Stuttgart, Baden-Württemberg, Germany.

- Remove vocal information (mostly done by html tags, so do it before you remove them)

Dr.-Ing. h.c. F. Porsche AG, usually shortened to Porsche (German pronunciation: [ˈpoʃə] (audio speaker iconlisten); see below), is a German automobile manufacturer specializing in high-performance sports cars, SUVs and sedans, headquartered in Stuttgart, Baden-Württemberg, Germany.

- Remove special characters

Dr.-Ing. h.c. F. Porsche AG, usually shortened to Porsche (:();), is a German automobile manufacturer specializing in high-performance sports cars, SUVs and sedans, headquartered in Stuttgart, Baden-Württemberg, Germany.

9.1 Text Processing Example III

- Convert to same case

```
Dr Ing h c F Porsche AG usually shortened to Porsche German is a German automobile manufacturer specializing in high-performance sports cars SUVs and sedans headquartered in Stuttgart Baden-Württemberg Germany
```

- Remove single characters, and unimportant words

```
dr ing h c f porsche ag usually shortened to porsche german is a german automobile manufacturer specializing in high-performance sports cars suvs and sedans headquartered in stuttgart baden-württemberg germany
```

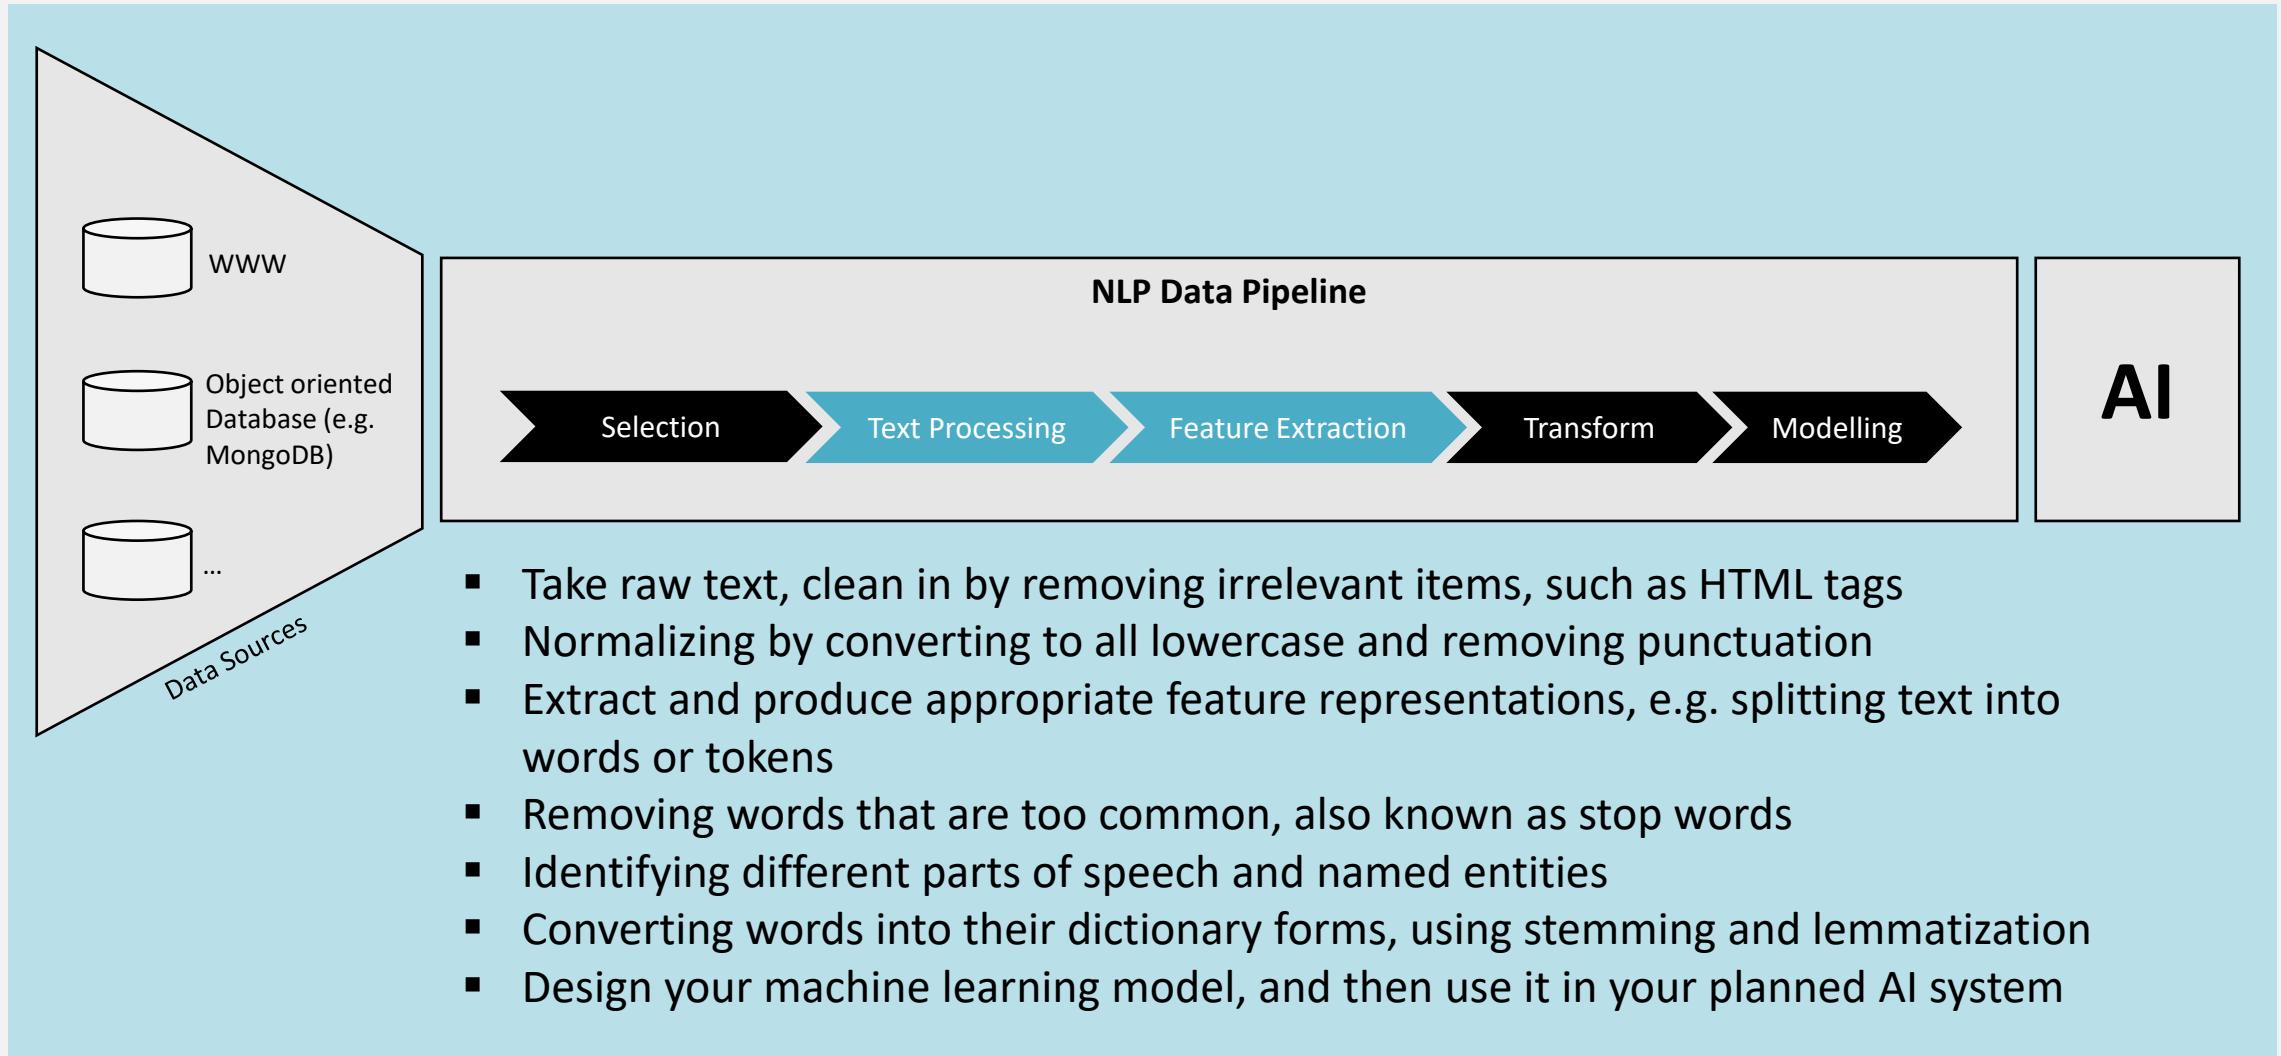
- Remove special characters

```
dr ing porsche ag shortened porsche german german automobile manufacturer specializing in high-performance sports cars suvs sedans headquartered stuttgart baden-württemberg germany
```



Transform to usable schemata like e.g. table

9.1 NLP Data Pipelines



9.1 NLP Cleaning Tools

Beautiful Soup

Beautiful Soup 4.9.0 documentation » Beautiful Soup Documentation

Table of Contents

- Beautiful Soup Documentation
 - Getting help
 - Quick Start
 - Installing Beautiful Soup
 - Installing a parser
 - Making the soup
 - Kinds of objects
 - Tag
 - Name
 - Attributes
 - Multi-valued attributes
 - NavigableString
 - BeautifulSoup
 - Comments and other special strings
 - Navigating the tree
 - Going down
 - Navigating using tag names
 - .contents and .children
 - .descendants
 - .string
 - .strings and .stripped_strings
 - Going up
 - parent
 - parents
 - Going sideways
 - .next_sibling and .previous_sibling
 - .next_siblings and .previous_siblings
 - Going back and forth
 - .next_element and .previous_element
 - .next_elements and .previous_elements
 - Searching the tree
 - Kinds of filters
 - A string
 - A regular expression
 - A list
 - True
 - A function

Beautiful Soup Documentation

Beautiful Soup is a Python library for pulling data out of HTML and XML files. It works with your favorite parser to provide idiomatic ways of navigating, searching, and modifying the parse tree. It commonly saves programmers hours or days of work.



These instructions illustrate all major features of Beautiful Soup 4, with examples. I show you what the library is good for, how it works, how to use it, how to make it do what you want, and what to do when it violates your expectations.

This document covers Beautiful Soup version 4.10.0. The examples in this documentation were written for Python 3.8.

You might be looking for the documentation for Beautiful Soup 3. If so, you should know that Beautiful Soup 3 is no longer being developed and that all support for it was dropped on December 31, 2020. If you want to learn about the differences between Beautiful Soup 3 and Beautiful Soup 4, see [Porting code to BS4](#).

This documentation has been translated into other languages by Beautiful Soup users:

- 这篇文档当然还有中文版。
- このページは日本語で利用できます(外部リンク)
- 이 문서는 한국어 번역도 가능합니다.
- Este documento também está disponível em Português do Brasil.
- Эта документация доступна на русском языке.

Getting help

If you have questions about Beautiful Soup, or run into problems, send mail to the discussion group. If your problem involves parsing an HTML document, be sure to mention what the `diagnose()` function says about that document.

Quick Start

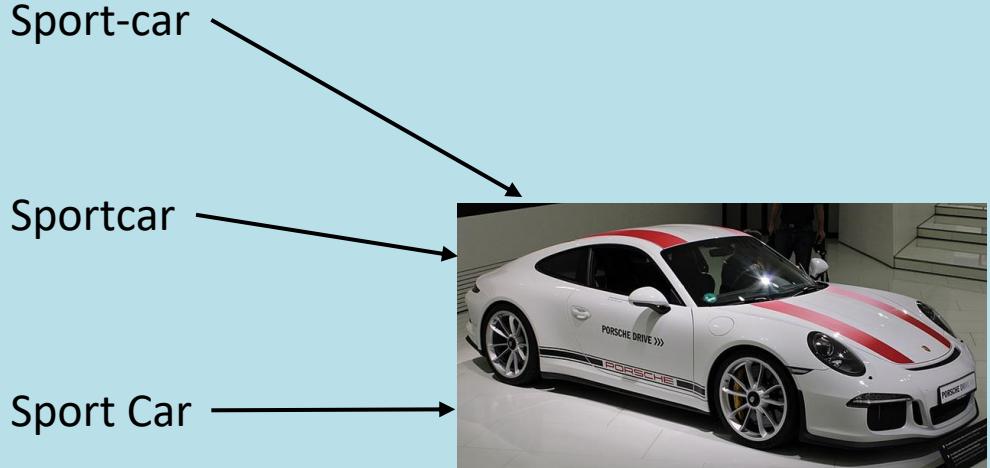
Here's an HTML document I'll be using as an example throughout this document. It's part of a story from *Alice in Wonderland*:

```
html_doc = """<html><head><title>The Dormouse's story</title></head>
<body>
```

Regular Expressions

```
/ ((\d{3}) (?:\.\|-)) ? (\d{3}) (?:\.\|-) (\d{4}) /
```

9.1 NLP Normalization



Sport-car

Sportcar

Sport Car

sport-car

...

```
text = "sport-Car"  
  
# convert to lowercase  
text = text.lower()  
  
# Remove punctuation characters  
import re  
text = re.sub(r"[^a-zA-Z0-9]", " ", text)
```

```
>>> print(text)
```



Do not mix up this concept with statistical normalization we discussed in lecture 4

Image Source: [Porsche 911 R im Porsche Museum 2018](#) (2018) by [Alexander Migl](#) from [Wikimedia](#) / [CC-BY-SA-4.0](#) (image not edited)

9.1 NLTK – Python package for Tokenization

NLTK

Documentation

```
import nltk  
nltk.download()
```

Search

nltk.tokenize package

NLTK Documentation

API Reference

Example Usage

Module Index

Wiki

FAQ

Open Issues

NLTK on GitHub

Installation

Installing NLTK

Installing NLTK Data

More

Release Notes

Contributing to NLTK

NLTK Team

NLTK Tokenizer Package

Tokenizers divide strings into lists of substrings. For example, tokenizers can be used to find the words and punctuation in a string:

```
>>> from nltk.tokenize import word_tokenize  
>>> s = '''Good muffins cost $3.88\nin New York. Please buy me  
... two of them.\n\nThanks.'''  
>>> word_tokenize(s)  
['Good', 'muffins', 'cost', '$', '3.88', 'in', 'New', 'York', '.',  
'Please', 'buy', 'me', 'two', 'of', 'them', '.', 'Thanks', '.']
```



If you want to save storage, you can download only specific functions like “punkt”, with `nltk.download('punkt')`. You can learn more about NLTK data installation on their ↗ [website](#)

This particular tokenizer requires the Punkt sentence tokenization models to be installed. NLTK also provides a simpler, regular-expression based tokenizer, which splits text on whitespace and punctuation:

```
>>> from nltk.tokenize import wordpunct_tokenize  
>>> wordpunct_tokenize(s)  
['Good', 'muffins', 'cost', '$', '3', '.', '88', 'in', 'New', 'York', '.',  
'Please', 'buy', 'me', 'two', 'of', 'them', '.', 'Thanks', '.']
```

9.1 Tokenization

- You can use the in-build tokenization methods from NLTK to perform straight forward tokenization of your texts

```
>>> from nltk.tokenize import word_tokenize  
>>> text = "I like this lecture. It has many practical examples."  
>>> # Split text into words using NLTK  
>>> words = word_tokenize(text)  
>>> print(words)  
  
['I', 'like', 'this', 'lecture', '.', 'It', 'has', 'many', 'practical',  
'examples', '.']
```



There are many more build-in tokenization functions. What will e.g.
`sent_tokenize(text)` return?

9.1 Stop Word Removal

- You can also use this package to remove noise like stopwords in different languages:

```
from nltk.corpus import stopwords  
words = [w for w in words if w not in stopwords.words("english")]
```

```
>>> print(stopwords.words("english"))  
['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you',  
"you're", "you've", "you'll", "you'd", 'your', 'yours', 'yourself',  
'yourselves', 'he', 'him', 'his', 'himself', 'she', "she's", 'her',  
'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them',  
'their', 'theirs ...
```

9.1 Motivation for Part-of-Speech Tagging (POS)

- Identifying how words are being used in a sentence helps us better understand what is being said:

Because she was so fast, Marie was always the first at the restaurant

pronoun verb noun

- Helps us to find relationships between words
- Get full list of word classes with `nltk.help.upenn_tagset()`

9.1 Part-of-Speech Tagging (POS)

- You can also use this package to remove noise like stopwords in different languages:

```
from nltk import pos_tag, ne_chunk
from nltk.tokenize import word_tokenize
text = "If ever there was proof that going electric needn't be a
compromise, the Porsche Taycan is it"
```

```
>>> sentence = word_tokenize(text) # tokenize text
>>> pos_tag(sentence) # tag each word with pos
```

[('If', 'IN'),	IN: preposition or conjunction, subordinating
('ever', 'RB'),	RB: adverb
('there', 'EX'),	EX: existential there
('was', 'VBD'),	VBD: verb, past tense
('proof', 'NN'),...	NN: noun, common, singular or mass

9.1 Named Entity Recognition (NER)

- Identifying how words are being used in a sentence helps us better understand what is being said:

2 words, 1 entity
The **Porsche AG** builds sport-cars

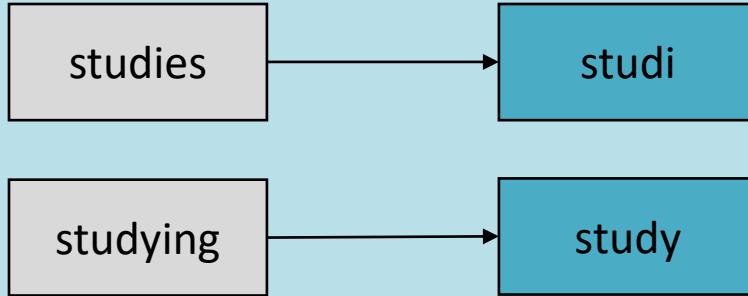
- Tokenize your text, make POS, then recognize named entities in text:

```
text = "The Porsche AG builds sport-cars"  
tree = ne_chunk(pos_tag(word_tokenize(text)))  
print(tree)
```

```
(S  
  The/DT  
  (ORGANIZATION Porsche/NNP)  
  AG/NNP  
  builds/VBZ  
  sport-cars/NNS)
```

9.1 Stemming and Lemmatization

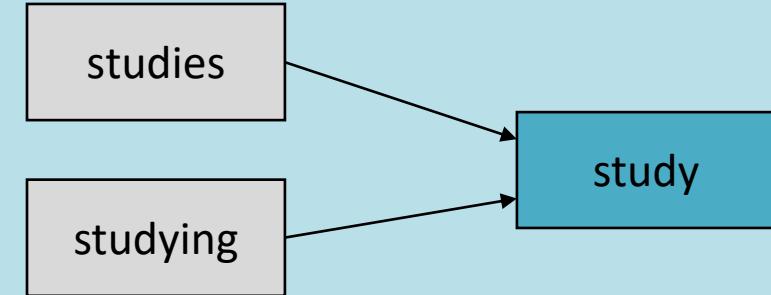
Stemming



- Removes suffixes

```
from nltk.stem import PorterStemmer  
from nltk.tokenize import word_tokenize  
  
ps = PorterStemmer()  
  
words = ["driver", "drives", "drive"]  
for w in words:  
    print(w, " : ", ps.stem(w))
```

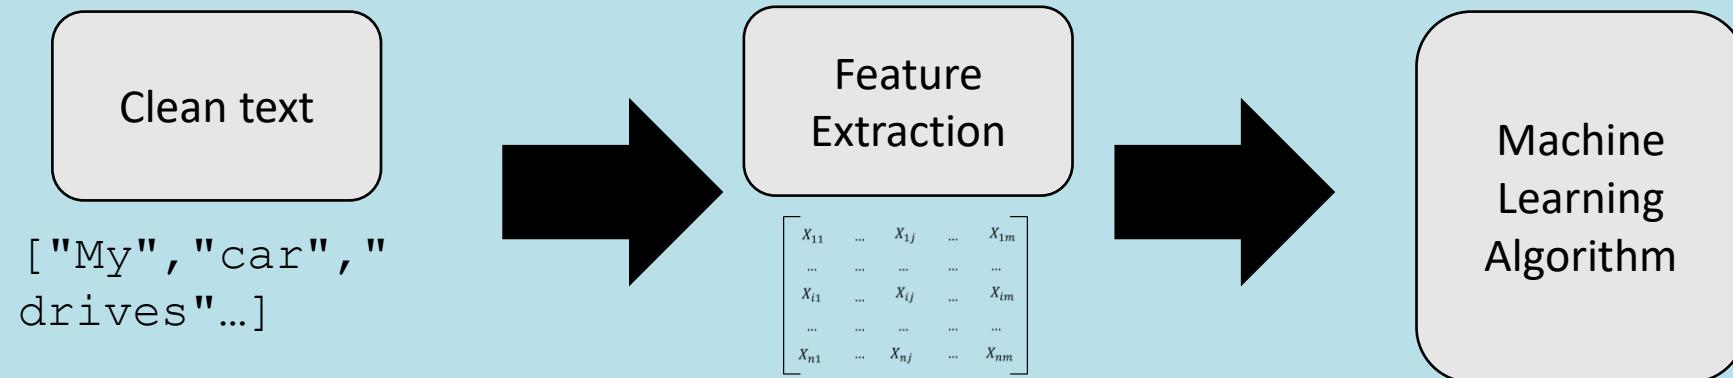
Lemmatization



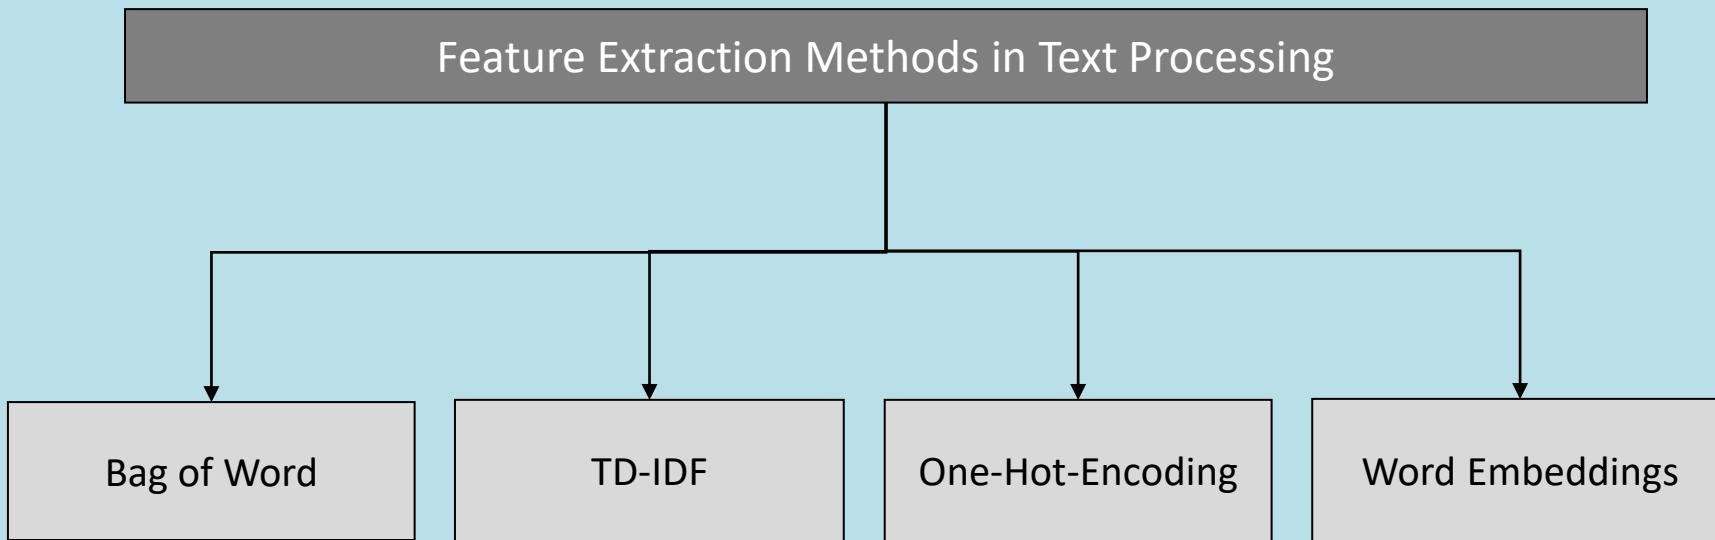
- Uses grammatical informations

```
from nltk.stem import WordNetLemmatizer  
  
lemmatizer = WordNetLemmatizer()  
  
print("sport-car : ",  
      lemmatizer.lemmatize("rocks"))
```

9.1 Motivation: Feature Extraction in Text Processing



9.1 Popular Word and Representations for NLP



9.1 Python Code Example for Feature Extraction

- Load Packages and tokenize documents

```
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer
from nltk.tokenize import word_tokenize

corpus = ["I like driving my Porsche.",
          "It is a very fast sport-car"
          ...]

stop_words = stopwords.words("english")
lemmatizer = WordNetLemmatizer()

def tokenize(text):
    text = re.sub(r"[^a-zA-Z0-9]", " ", text.lower())
    tokens = word_tokenize(text)
    tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]

    return tokens
```

9.1 Bag of Words and Document Frequency

```
Document1 =  
["My", "car",  
"drives",  
"very", "fast"]  
  
Document2 =  
["My", "car",  
"is", "a",  
"Porsche"]  
  
Document3 =  
["I", "like",  
"my", "911"]
```

Document-term matrix

Term	Doc1	Doc2	Doc3	Sum
car	1	1	0	2
drives	1	0	0	1
very	1	0	0	1
fast	1	0	0	1
Porsche	0	1	0	1
like	0	0	1	1
911	0	0	1	1

Noise = ["a", "is", "my", "I"]

9.1 Similarity in Bag of Words

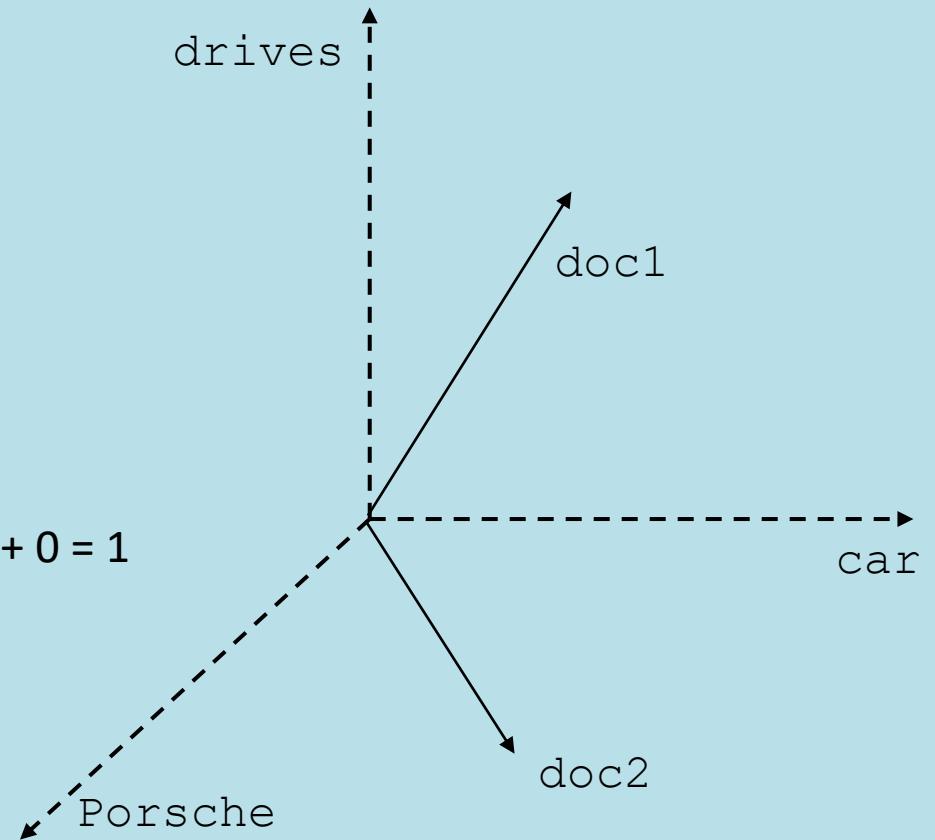
Term	Doc1	Doc2
car	1	1
drives	1	0
very	1	0
fast	1	0
Porsche	0	1

1. Dot product

$$doc1 \cdot doc2 = \sum a_o b_o + a_1 b_1 + \dots + a_n b_n = 1 + 0 + 0 + 0 + 0 = 1$$

2. Cosine similarity

$$\cos(\theta) = \frac{doc1 \cdot doc2}{\| doc1 \| \cdot \| doc2 \|} = \frac{1}{\sqrt{4} \times \sqrt{2}}$$



9.1 Bag of Words with NLTK

- initialize count vectorizer object

```
from sklearn.feature_extraction.text import CountVectorizer  
vect = CountVectorizer(tokenizer=tokenize)
```

- compute counts of each word (token) in our text data

```
>>> X = vect.fit_transform(corpus)  
>>> X.toarray()  
  
array([[0, 0, 0, 1, 0], [0, 0, 0, 1, 0], ...], dtype=int64)
```

9.1 Term Frequency (TF) – Inverse Document Frequency (IDF)

Term	Doc1	Doc2	Doc3	Freq
engine	3	2	0	5
drives	2	1	0	3
fuel	3	1	0	4
price	0	0	5	5
saldo	0	0	3	3
...

Term	Doc1	Doc2	Doc3	Freq
engine				5
drives				3
fuel				4
price				5
saldo				3
...

$$tf = \frac{count(t, d)}{|d|}$$

$$idf = \log\left(\frac{|D|}{|\{d \in D : t \in d\}|}\right)$$

$$tfidf = tf \cdot idf$$



What will be the new values of the idf table? Can you compute them?

Do you have any idea of the background of doc1, doc2 and doc3?

9.1 TfIdfTransformer with sklearn

- Compute tf-idf values based on the counts from count vectorizer

```
from sklearn.feature_extraction.text import TfidfTransformer  
  
transformer = TfidfTransformer(smooth_idf=False)  
tfidf = transformer.fit_transform(X)
```

- You can print it if you want

```
>>> tfidf.toarray()  
  
array([[ 0.        ,  0.        ,  0.        ,  0.        ,  0.37      ,  0.        ,  
         0.        ,  0.        ,  0.        ,  0.        ,  0.        ,  0.        ,  
         0.        ,  0.        ,  0.        ,  0.        ,  0.        ,  0.57      ,  0.  
         0.        ,  0.        ,  0.        ,  0.        ,  0.        ,  0.        ]])
```

9.1 Alternative: TfidfVectorizer = CountVectorizer + TfidfTransformer

- compute tf-idf values based on the counts from count vectorizer

```
from sklearn.feature_extraction.text import TfidfVectorizer

# initialize tf-idf vectorizer object
vectorizer = TfidfVectorizer()

# compute bag of word counts and tf-idf values
X = vectorizer.fit_transform(corpus)
```

9.1 Application of TF-IDF score

- **Information retrieval:** Match TF-IDF score of a user query against the whole document set to figure out how relevant a document is to that given query
- **Keywords extraction:** The highest ranking words for a document in terms of TF-IDF score can very well represent the keywords of that document

Your turn!

Task

Please explain

- What is Information Retrieval
- Name an application of Information Retrieval that we didn't discuss in lecture
- Explain the role of NLP in Information Retrieval

Outline

9 Language and Image Processing

9.1 Applied Natural Language Processing

9.2 Speech Processing and Communication

9.3 Perception and Image Processing

9.4 Generative Modelling

Lectorial 7: Building NLP and Machine Learning Pipelines for Webservices

► **What we will learn:**

- We will discuss how to make use of the copious knowledge that is expressed in natural language
- And how to process and prepare language data for machines to build state-of-the-art machine learning pipelines for language processing
- We learn new type of models, so called generative models, which allow to generate new instances based on the knowledge of your model



Image source: [↗ Pixabay](#) (2021) / [↗ CC0](#)

► **Duration:**

- 135 min + 90 (Lectorial)

► **Relevant for Exam:**

- 9.1, 9.4

This Software Engineer Designed a Chatbot to Chat With His Girlfriend While He's Busy at Work

NEWS 18



However, the girl eventually got suspicious over the speed she was receiving messages from her "boyfriend."

FORTUNE

SEARCH SIGN IN Subscribe

Artificial Intelligence | Cryptocurrency | Metaverse | Cybersecurity | Tech Forward

TECH · ASHLEY MADISON

Ashley Madison Used Chatbots to Lure Cheaters, Then Threatened to Expose Them When They Complained

BY DAVID Z. MORRIS

July 10, 2016 at 7:11 PM GMT+2



9.2 Communication

Speaker

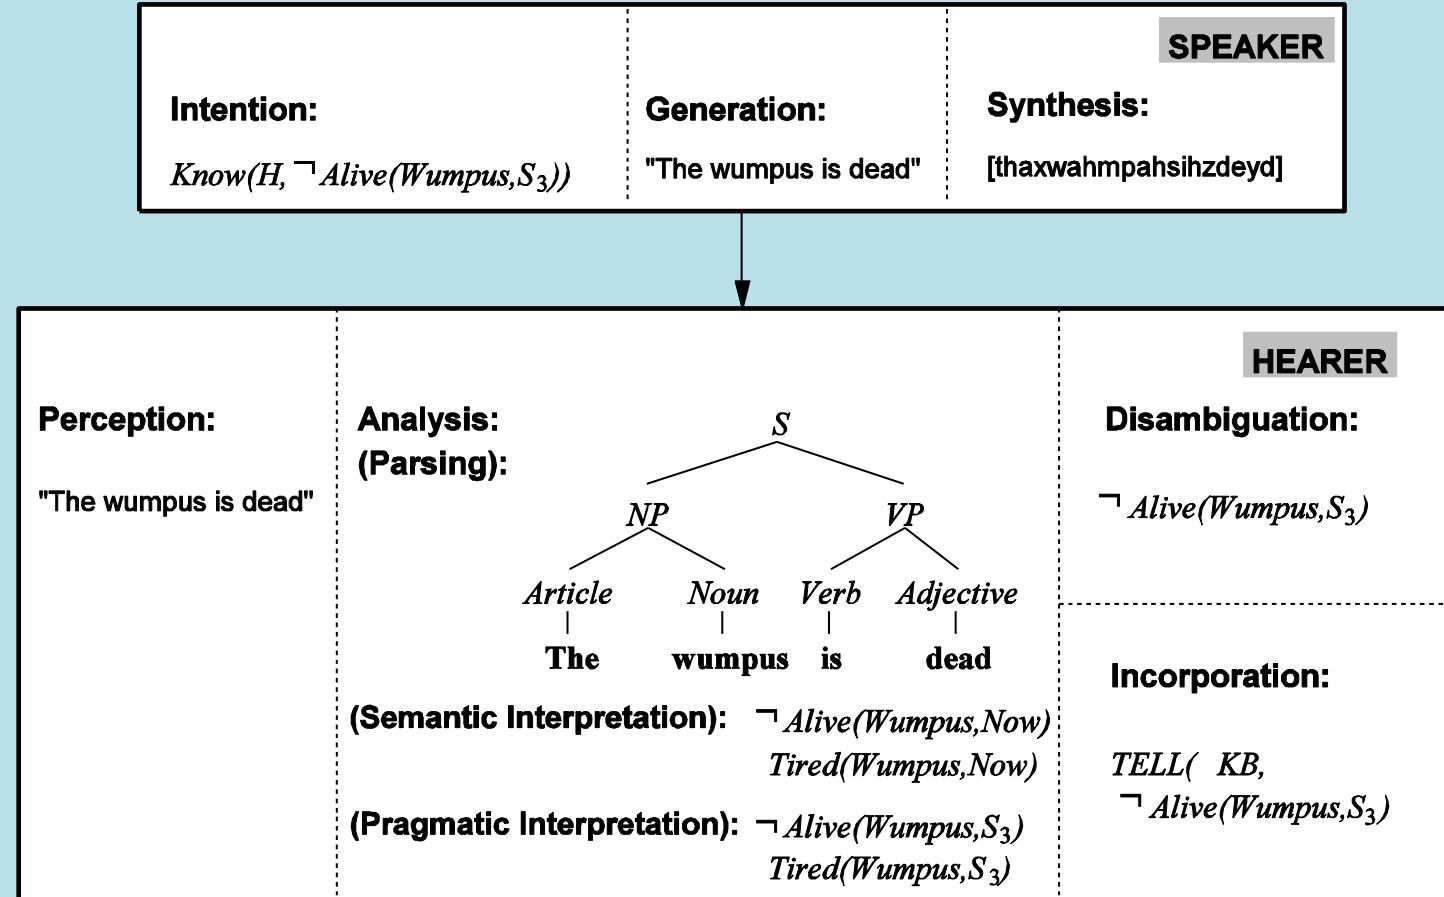
- **Intention:** Decide when and what information should be transmitted (a.k.a. strategic generation). May require planning and reasoning about agents' goals and beliefs.
- **Generation:** Translate the information to be communicated (in internal logical representation or “language of thought”) into string of words in desired natural language (a.k.a. tactical generation).
- **Synthesis:** Output the string in desired modality, text or speech.

Supervised Learning

- **Perception:** Map input modality to a string of words, e.g. optical character recognition (OCR) or speech recognition.
- **Analysis:** Determine the information content of the string.
 - Syntactic interpretation (parsing): Find the correct parse tree showing the phrase structure of the string.
 - Semantic Interpretation: Extract the (literal) meaning of the string (logical form).
 - Pragmatic Interpretation: Consider effect of the overall context on altering the literal meaning of a sentence.
- **Incorporation:** Decide whether or not to believe the content of the string and add it to the KB.

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Communication (cont)



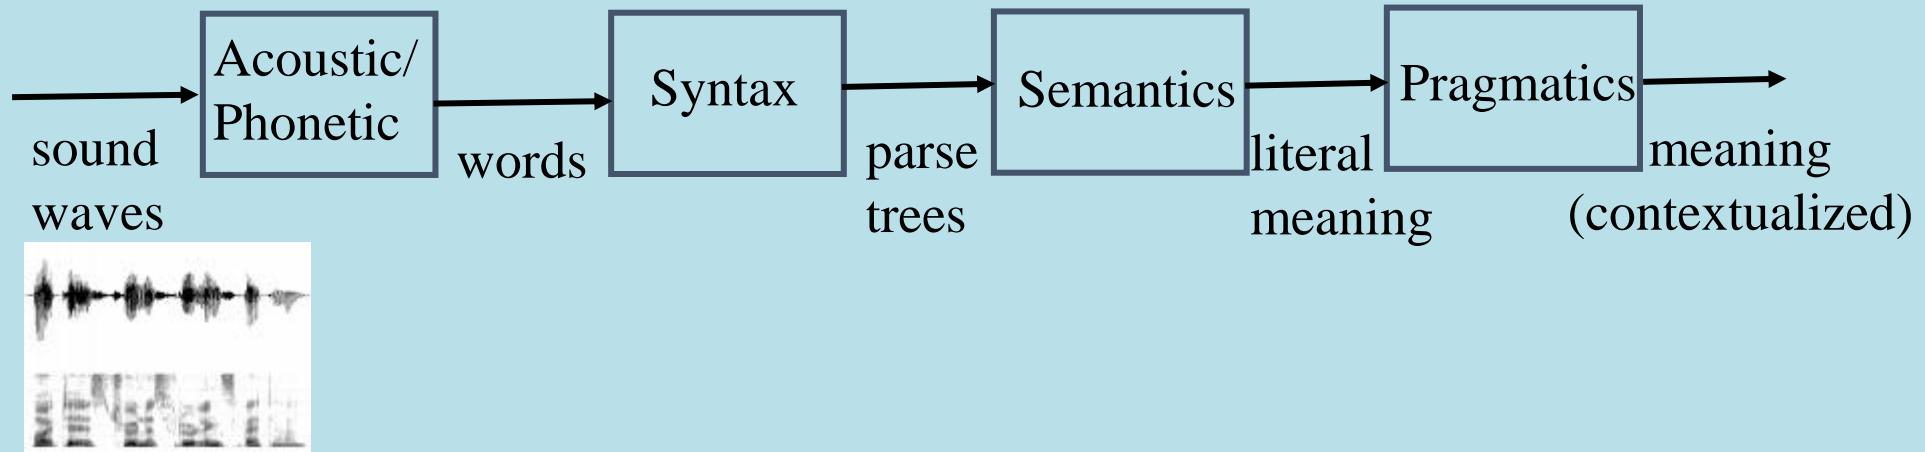
Adapted from Russell, S., & Norvig, P. (2016)

9.2 Syntax, Semantic, Pragmatics

- Syntax concerns the proper ordering of words and its affect on meaning.
 - The dog bit the boy.
 - The boy bit the dog.
 - Colorless green ideas sleep furiously.
- Semantics concerns the (literal) meaning of words, phrases, and sentences.
 - “plant” as a photosynthetic organism
 - “plant” as a manufacturing facility
 - “plant” as the act of sowing
- Pragmatics concerns the overall communicative and social context and its effect on interpretation.
 - The ham sandwich wants another beer. (co-reference, anaphora)
 - John thinks vanilla. (ellipsis)

Adapted from Russell, S., & Norvig, P. (2016)

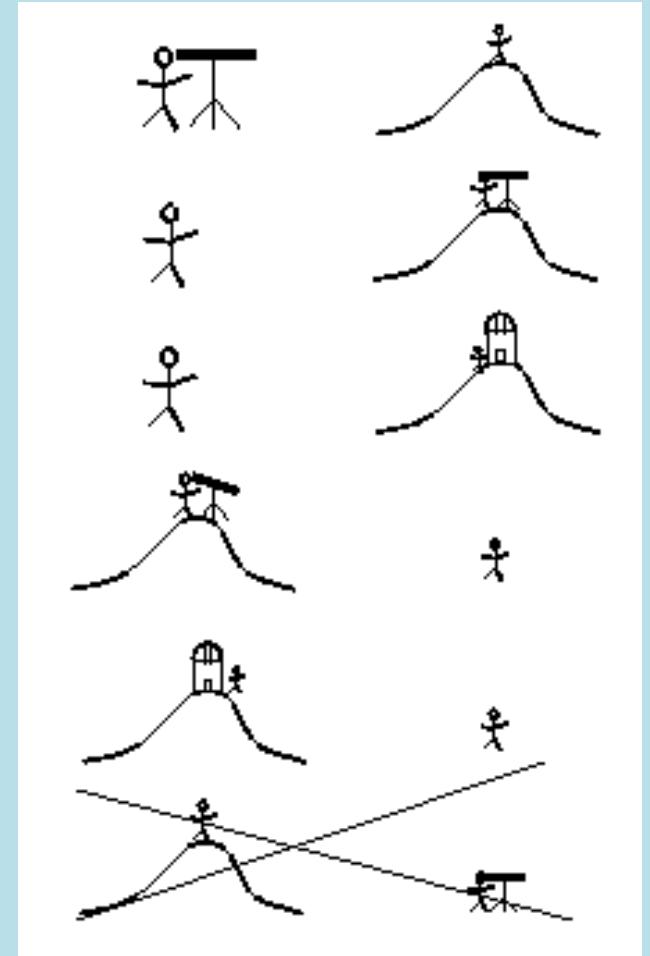
9.2 Modular Comprehension



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Ambiguity

- Natural language is highly ambiguous and must be *disambiguated*.
 - I saw the man on the hill with a telescope.
 - I saw the Grand Canyon flying to LA.
 - Time flies like an arrow.
 - Horse flies like a sugar cube.
 - Time runners like a coach.
 - Time cars like a Porsche.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Ambiguity is Ubiquitous I

- Speech Recognition

- “recognize speech” vs. “wreck a nice beach”
- “youth in Asia” vs. “euthanasia”

- Syntactic Analysis

- “I ate spaghetti with chopsticks” vs. “I ate spaghetti with meatballs.”

- Semantic Analysis

- “The dog is in the pen.” vs. “The ink is in the pen.”
- “I put the plant in the window” vs. “Ford put the plant in Mexico”

Adapted from Russell, S., & Norvig, P. (2016)

- Pragmatic Analysis
 - From “The Pink Panther Strikes Again”:

Clouseau: Does your dog bite?

Hotel Clerk: No.

Clouseau: [bowing down to pet the dog] Nice doggie.

[Dog barks and bites Clouseau in the hand]

Clouseau: I thought you said your dog did not bite!

Hotel Clerk: That is not my dog.

9.2 Ambiguity is Explosive

- Ambiguities compound to generate enormous numbers of possible interpretations.
- In English, a sentence ending in n prepositional phrases has over 2^n syntactic interpretations (cf. Catalan numbers).
 - “I saw the man with the telescope”: 2 parses
 - “I saw the man on the hill with the telescope.”: 5 parses
 - “I saw the man on the hill in Texas with the telescope”: 14 parses
 - “I saw the man on the hill in Texas with the telescope at noon.”: 42 parses
 - “I saw the man on the hill in Texas with the telescope at noon on Monday” 132 parses

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Humor and Ambiguity

Many jokes rely on the ambiguity of language:

- Groucho Marx: One morning I shot an elephant in my pajamas. How he got into my pajamas, I'll never know.
- She criticized my apartment, so I knocked her flat.
- Noah took all of the animals on the ark in pairs. Except the worms, they came in apples.
- Policeman to little boy: “We are looking for a thief with a bicycle.”
Little boy: “Wouldn’t you be better using your eyes.”
- Why is the teacher wearing sun-glasses. Because the class is so bright.

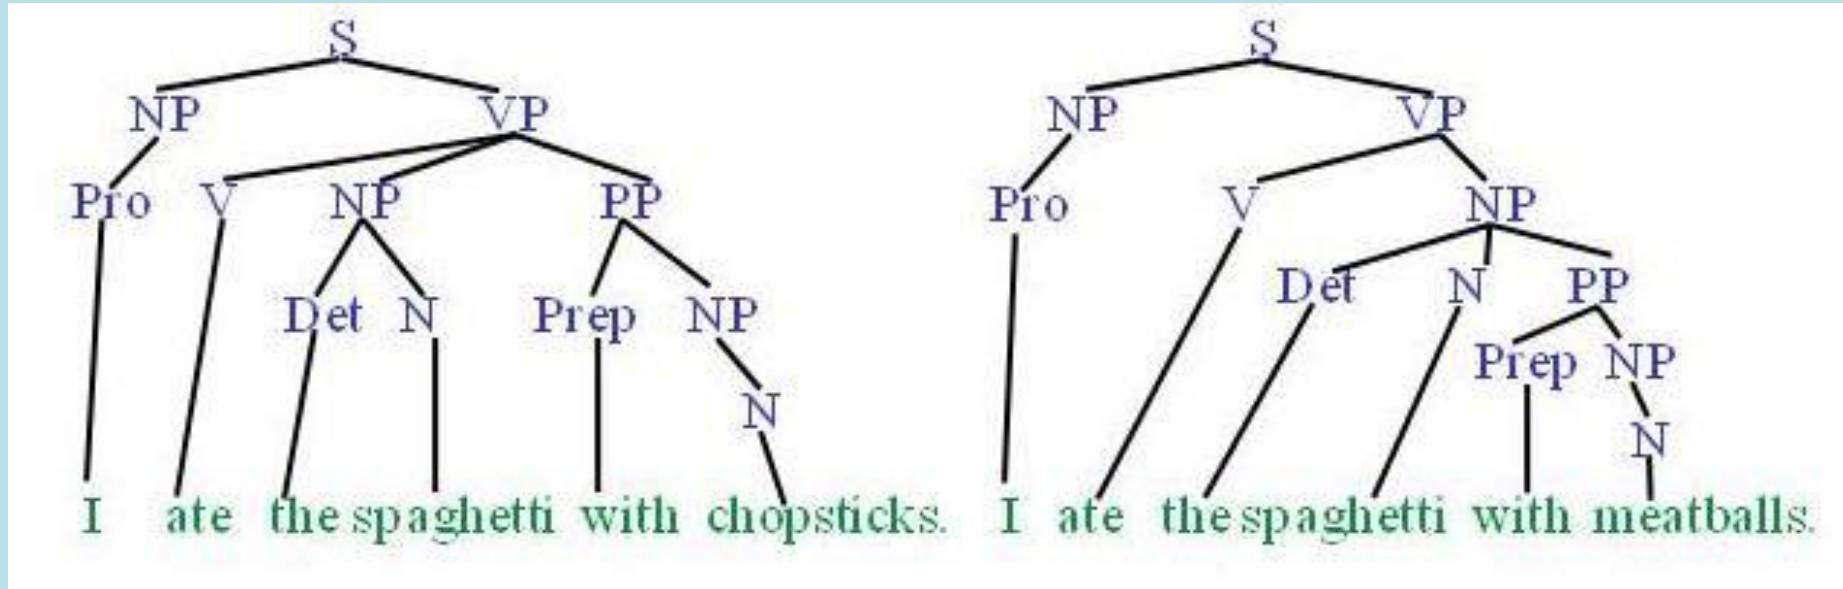
9.2 Natural Languages vs. Computer Languages

- Ambiguity is the primary difference between natural and computer languages.
- Formal programming languages are designed to be unambiguous, i.e. they can be defined by a grammar that produces a unique parse for each sentence in the language.
- Programming languages are also designed for efficient (deterministic) parsing, i.e. they are deterministic context-free languages (DCFLs).
- A sentence in a DCFL can be parsed in $O(n)$ time where n is the length of the string.

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Syntactic Parsing

- Produce the correct syntactic parse tree for a sentence.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Context Free Grammars (CFG)

- N a set of ***non-terminal symbols*** (or ***variables***)
- Σ a set of ***terminal symbols*** (disjoint from N)
- R a set of ***productions*** or ***rules*** of the form $A \rightarrow \beta$, where A is a non-terminal and β is a string of symbols from $(\Sigma \cup N)^*$
- S , a designated non-terminal called the ***start symbol***

9.2 Simple CFG for ATIS English

Grammar

$S \rightarrow NP\ VP$

$S \rightarrow Aux\ NP\ VP$

$S \rightarrow VP$

$NP \rightarrow Pronoun$

$NP \rightarrow Proper-Noun$

$NP \rightarrow Det\ Nominal$

$Nominal \rightarrow Noun$

$Nominal \rightarrow Nominal\ Noun$

$Nominal \rightarrow Nominal\ PP$

$VP \rightarrow Verb$

$VP \rightarrow Verb\ NP$

$VP \rightarrow VP\ PP$

$PP \rightarrow Prep\ NP$

Lexicon

$Det \rightarrow the | a | that | this$

$Noun \rightarrow book | flight | meal | money$

$Verb \rightarrow book | include | prefer$

$Pronoun \rightarrow I | he | she | me$

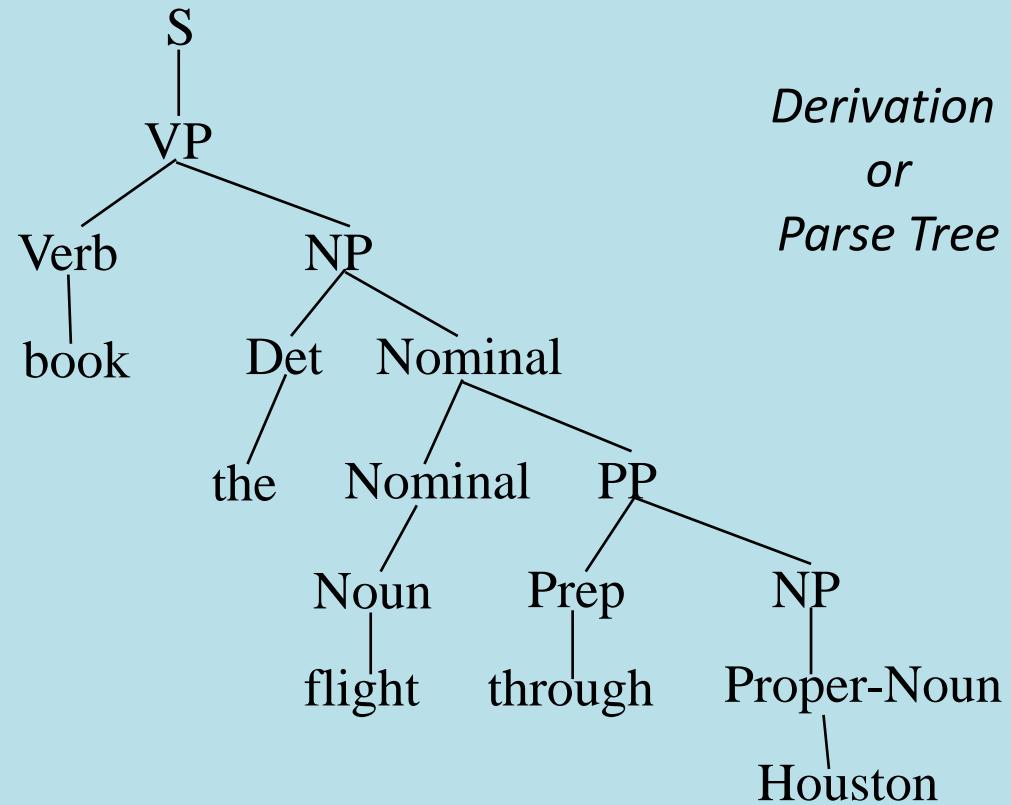
$Proper-Noun \rightarrow Houston | NWA$

$Aux \rightarrow does$

$Prep \rightarrow from | to | on | near | through$

Adapted from Russell, S., & Norvig, P. (2016)

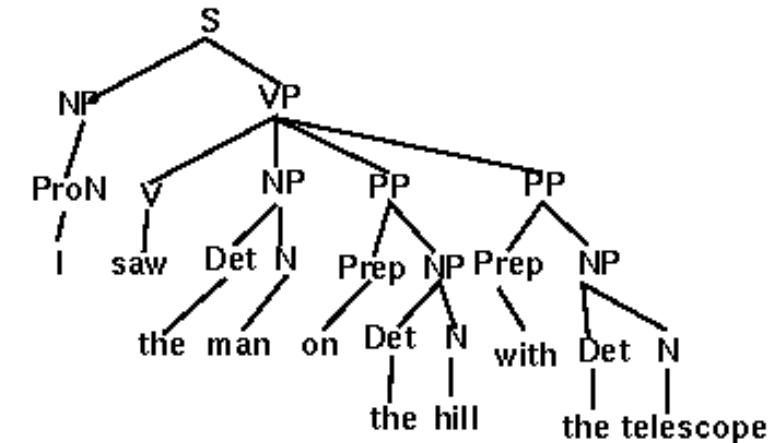
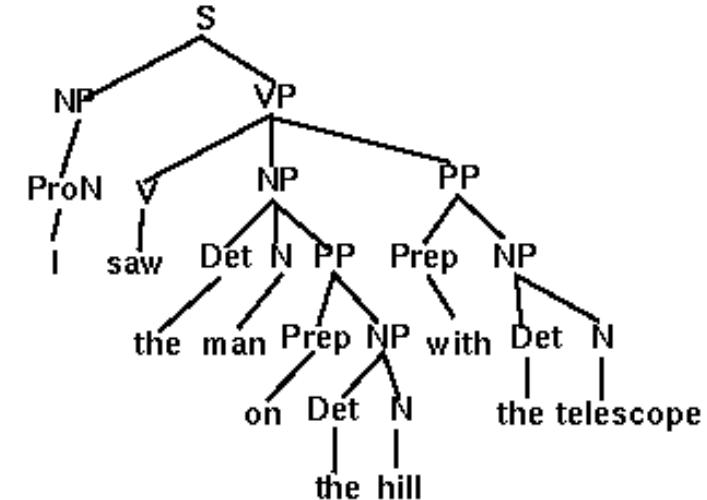
- Sentences are generated by recursively rewriting the start symbol using the productions until only terminals symbols remain.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Parse Trees and Syntactic Ambiguity

- If a sentence has more than one possible derivation (parse tree) it is said to be *syntactically ambiguous*.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Prepositional Phrase Attachment Explosion

- A transitive English sentence ending in m prepositional phrases has *at least* 2^m parses.

I saw the man on the hill with a telescope on Tuesday in Austin....

The diagram shows a sentence "I saw the man on the hill with a telescope on Tuesday in Austin...." with several curved arrows originating from the verb "saw". One arrow points to "the man", another to "on the hill", one to "with a telescope", and two to "on Tuesday". This illustrates the multiple ways a single verb can be attached to different prepositional phrases.

- The exact number of parses is given by the *Catalan numbers* (where $n=m+1$)

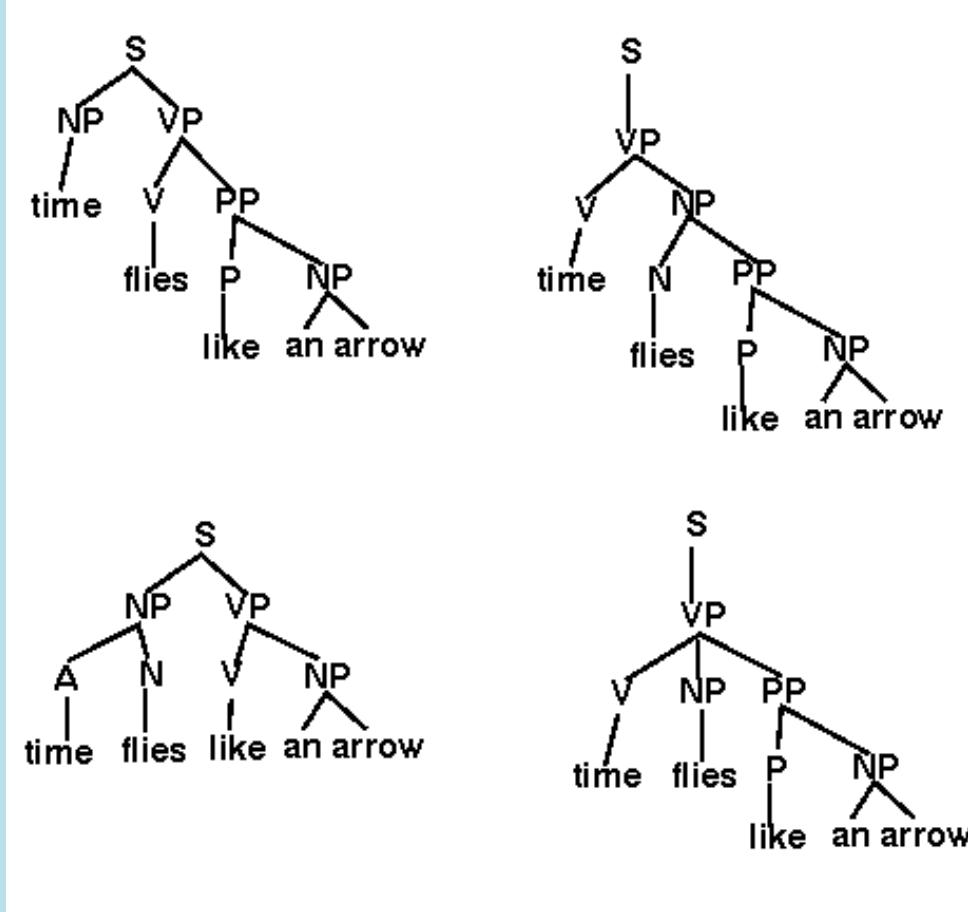
$$\binom{2n}{n} - \binom{2n}{n-1} \approx \frac{4^n}{n^{3/2} \sqrt{\pi}}$$

1, 2, 5, 14, 132, 429, 1430, 4862, 16796,

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Spurious Ambiguity

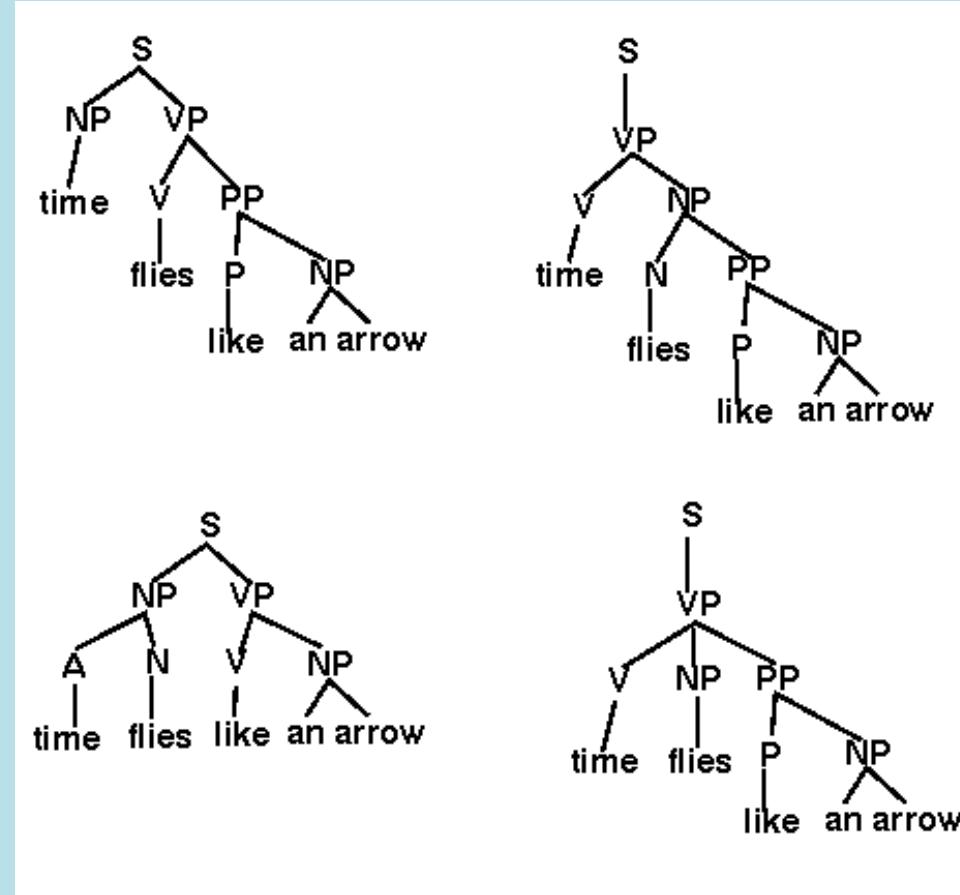
- Most parse trees of most NL sentences make no sense.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Spurious Ambiguity

- Most parse trees of most NL sentences make no sense.



Adapted from Russell, S., & Norvig, P. (2016)

- Given a string of non-terminals and a CFG, determine if the string can be generated by the CFG.
 - Also return a parse tree for the string
 - Also return all possible parse trees for the string
- Must search space of derivations for one that derives the given string.
 - **Top-Down Parsing:** Start searching space of derivations for the start symbol.
 - **Bottom-up Parsing:** Start search space of reverse deviations from the terminal symbols in the string.

9.2 Top Down vs. Bottom Up

- Top down never explores options that will not lead to a full parse, but can explore many options that never connect to the actual sentence.
- Bottom up never explores options that do not connect to the actual sentence but can explore options that can never lead to a full parse.
- Relative amounts of wasted search depend on how much the grammar branches in each direction.

9.2 Syntactic Parsing & Ambiguity

- Just produces all possible parse trees.
- Does not address the important issue of ambiguity resolution.

9.2 Statistical Parsing

- Statistical parsing uses a probabilistic model of syntax in order to assign probabilities to each parse tree.
- Provides principled approach to resolving syntactic ambiguity.
- Allows supervised learning of parsers from tree-banks of parse trees provided by human linguists.
- Also allows unsupervised learning of parsers from unannotated text, but the accuracy of such parsers has been limited.

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Probabilistic Context Free Grammar (PCFG)

- A PCFG is a probabilistic version of a CFG where each production has a probability.
- Probabilities of all productions rewriting a given non-terminal must add to 1, defining a distribution for each non-terminal.
- String generation is now probabilistic where production probabilities are used to non-deterministically select a production for rewriting a given non-terminal.

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Simple PCFG for ATIS English

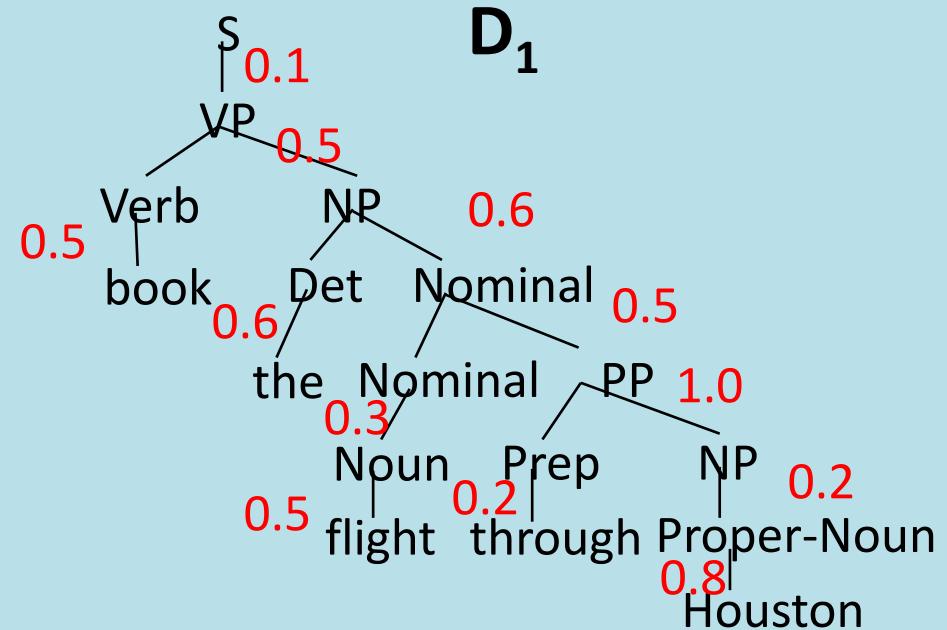
Grammar	Prob	Lexicon
$S \rightarrow NP\ VP$	0.8	Det → the a that this 0.6 0.2 0.1 0.1
$S \rightarrow Aux\ NP\ VP$	0.1 + 1.0	Noun → book flight meal money 0.1 0.5 0.2 0.2
$S \rightarrow VP$	0.1	
$NP \rightarrow Pronoun$	0.2	
$NP \rightarrow Proper-Noun$	0.2 + 1.0	Verb → book include prefer 0.5 0.2 0.3
$NP \rightarrow Det\ Nominal$	0.6	
$Nominal \rightarrow Noun$	0.3	Pronoun → I he she me 0.5 0.1 0.1 0.3
$Nominal \rightarrow Nominal\ Noun$	0.2 + 1.0	Proper-Noun → Houston NWA 0.8 0.2
$Nominal \rightarrow Nominal\ PP$	0.5	
$VP \rightarrow Verb$	0.2	
$VP \rightarrow Verb\ NP$	0.5 + 1.0	Aux → does 1.0
$VP \rightarrow VP\ PP$	0.3	
$PP \rightarrow Prep\ NP$	1.0	Prep → from to on near through 0.25 0.25 0.1 0.2 0.2

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Sentence Probability

- Assume productions for each node are chosen independently.
- Probability of derivation is the product of the probabilities of its productions.

$$\begin{aligned} P(D_1) &= 0.1 \times 0.5 \times 0.5 \times 0.6 \times 0.6 \times \\ &\quad 0.5 \times 0.3 \times 1.0 \times 0.2 \times 0.2 \times \\ &\quad 0.5 \times 0.8 \\ &= 0.0000216 \end{aligned}$$

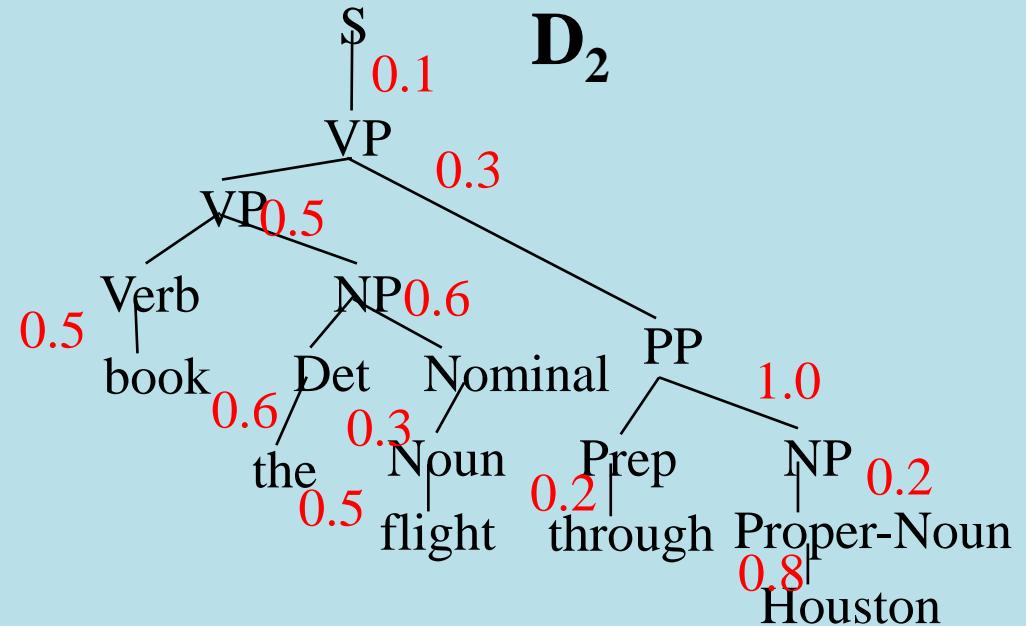


Adapted from Russell, S., & Norvig, P. (2016)

9.2 Syntactic Disambiguation

- Resolve ambiguity by picking most probable parse tree.

$$\begin{aligned} P(D_2) &= 0.1 \times 0.3 \times 0.5 \times 0.6 \times 0.5 \times \\ &\quad 0.6 \times 0.3 \times 1.0 \times 0.5 \times 0.2 \times \\ &\quad 0.2 \times 0.8 \\ &= 0.00001296 \end{aligned}$$



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Sentence Probability

- Probability of a sentence is the sum of the probabilities of all of its derivations.

$$\begin{aligned} P(\text{"book the flight through Houston"}) &= \\ P(D_1) + P(D_2) &= 0.0000216 + 0.00001296 \\ &= 0.00003456 \end{aligned}$$

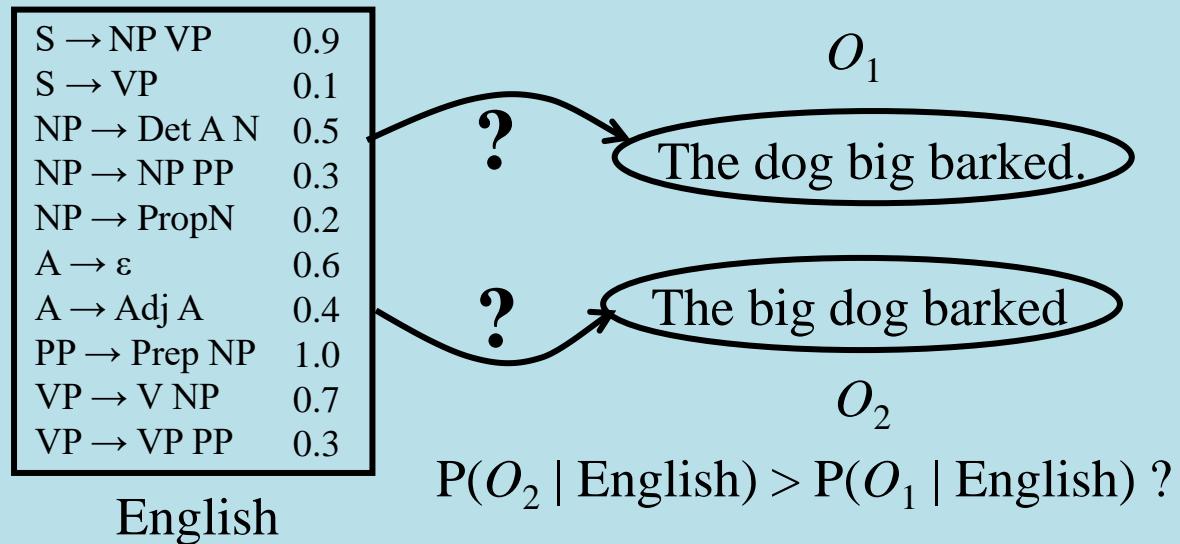
Adapted from Russell, S., & Norvig, P. (2016)

9.2 Three Useful PCFG Tasks

- Observation likelihood: To classify and order sentences.
- Most likely derivation: To determine the most likely parse tree for a sentence.
- Maximum likelihood training: To train a PCFG to fit empirical training data.

9.2 PCFG: Observation Likelihood

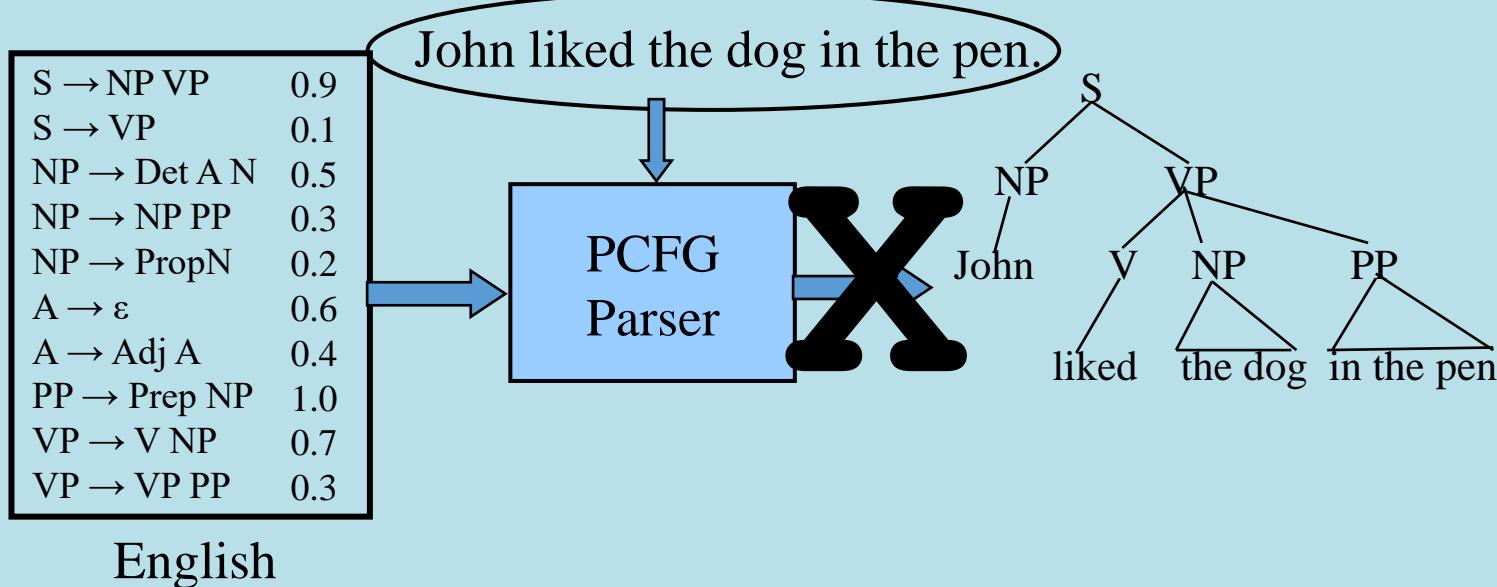
- What is the probability that a given string is produced by a given PCFG
- Can use a PCFG as a language model to choose between alternative sentences for speech recognition or machine translation



Adapted from Russell, S., & Norvig, P. (2016)

9.2 PCFG: Most Likely Derivation

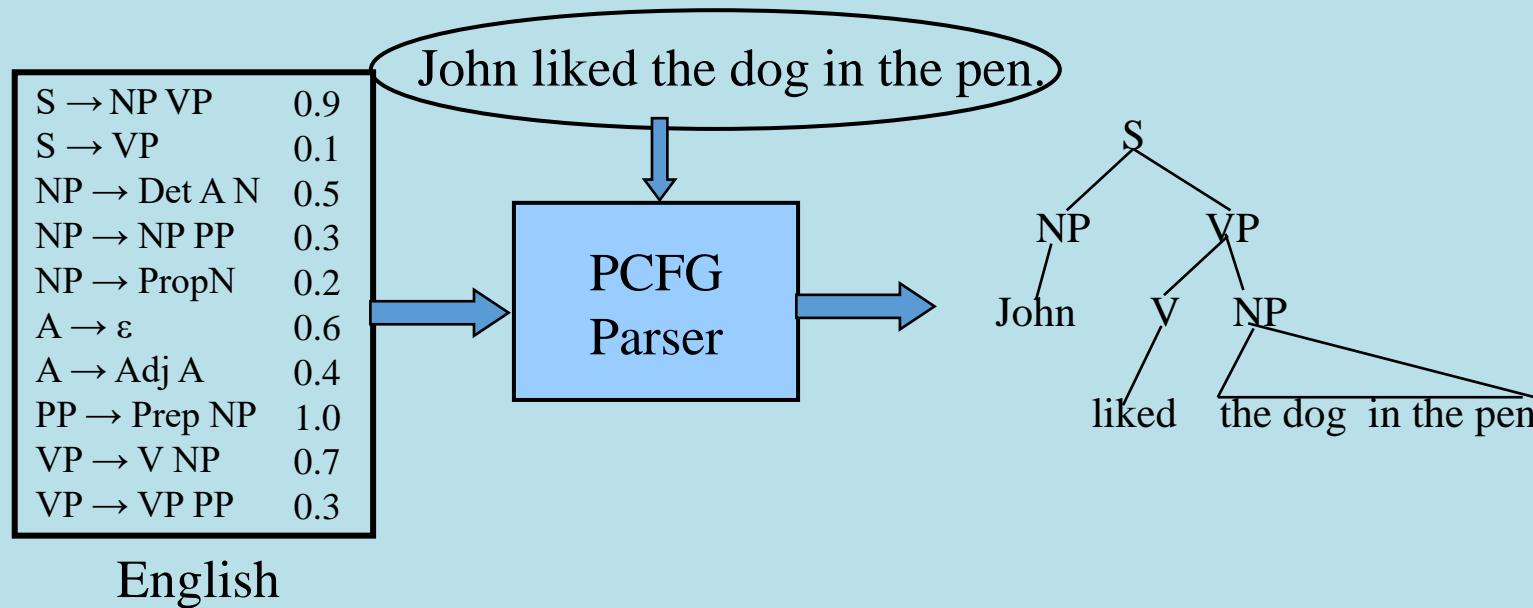
- What is the most probable derivation (parse tree) for a sentence.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 PCFG: Most Likely Derivation

- What is the most probable derivation (parse tree) for a sentence.

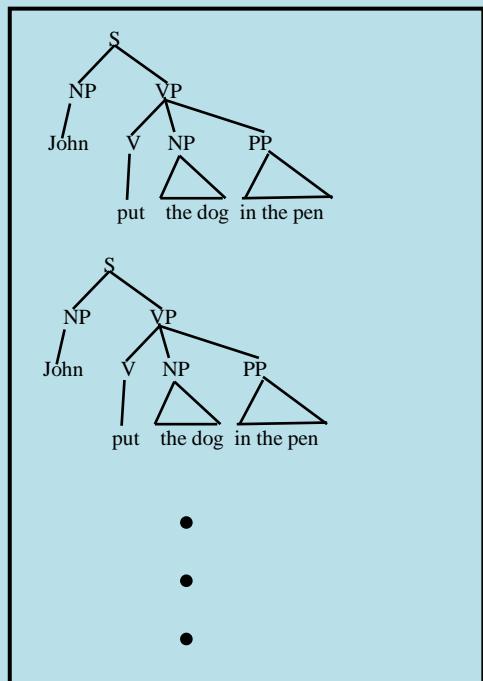


Adapted from Russell, S., & Norvig, P. (2016)

9.2 PCFG: Supervised Training

- If parse trees are provided for training sentences, a grammar and its parameters can all be estimated directly from counts accumulated from the tree-bank (with appropriate smoothing).

Tree Bank



Supervised
PCFG
Training

$S \rightarrow NP\ VP$	0.9
$S \rightarrow VP$	0.1
$NP \rightarrow Det\ A\ N$	0.5
$NP \rightarrow NP\ PP$	0.3
$NP \rightarrow PropN$	0.2
$A \rightarrow \epsilon$	0.6
$A \rightarrow Adj\ A$	0.4
$PP \rightarrow Prep\ NP$	1.0
$VP \rightarrow V\ NP$	0.7
$VP \rightarrow VP\ PP$	0.3

English

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Estimating Production Probabilities

- Set of production rules can be taken directly from the set of rewrites in the treebank.
- Parameters can be directly estimated from frequency counts in the treebank.

$$P(\alpha \rightarrow \beta | \alpha) = \frac{\text{count}(\alpha \rightarrow \beta)}{\sum_{\gamma} \text{count}(\alpha \rightarrow \gamma)} = \frac{\text{count}(\alpha \rightarrow \beta)}{\text{count}(\alpha)}$$

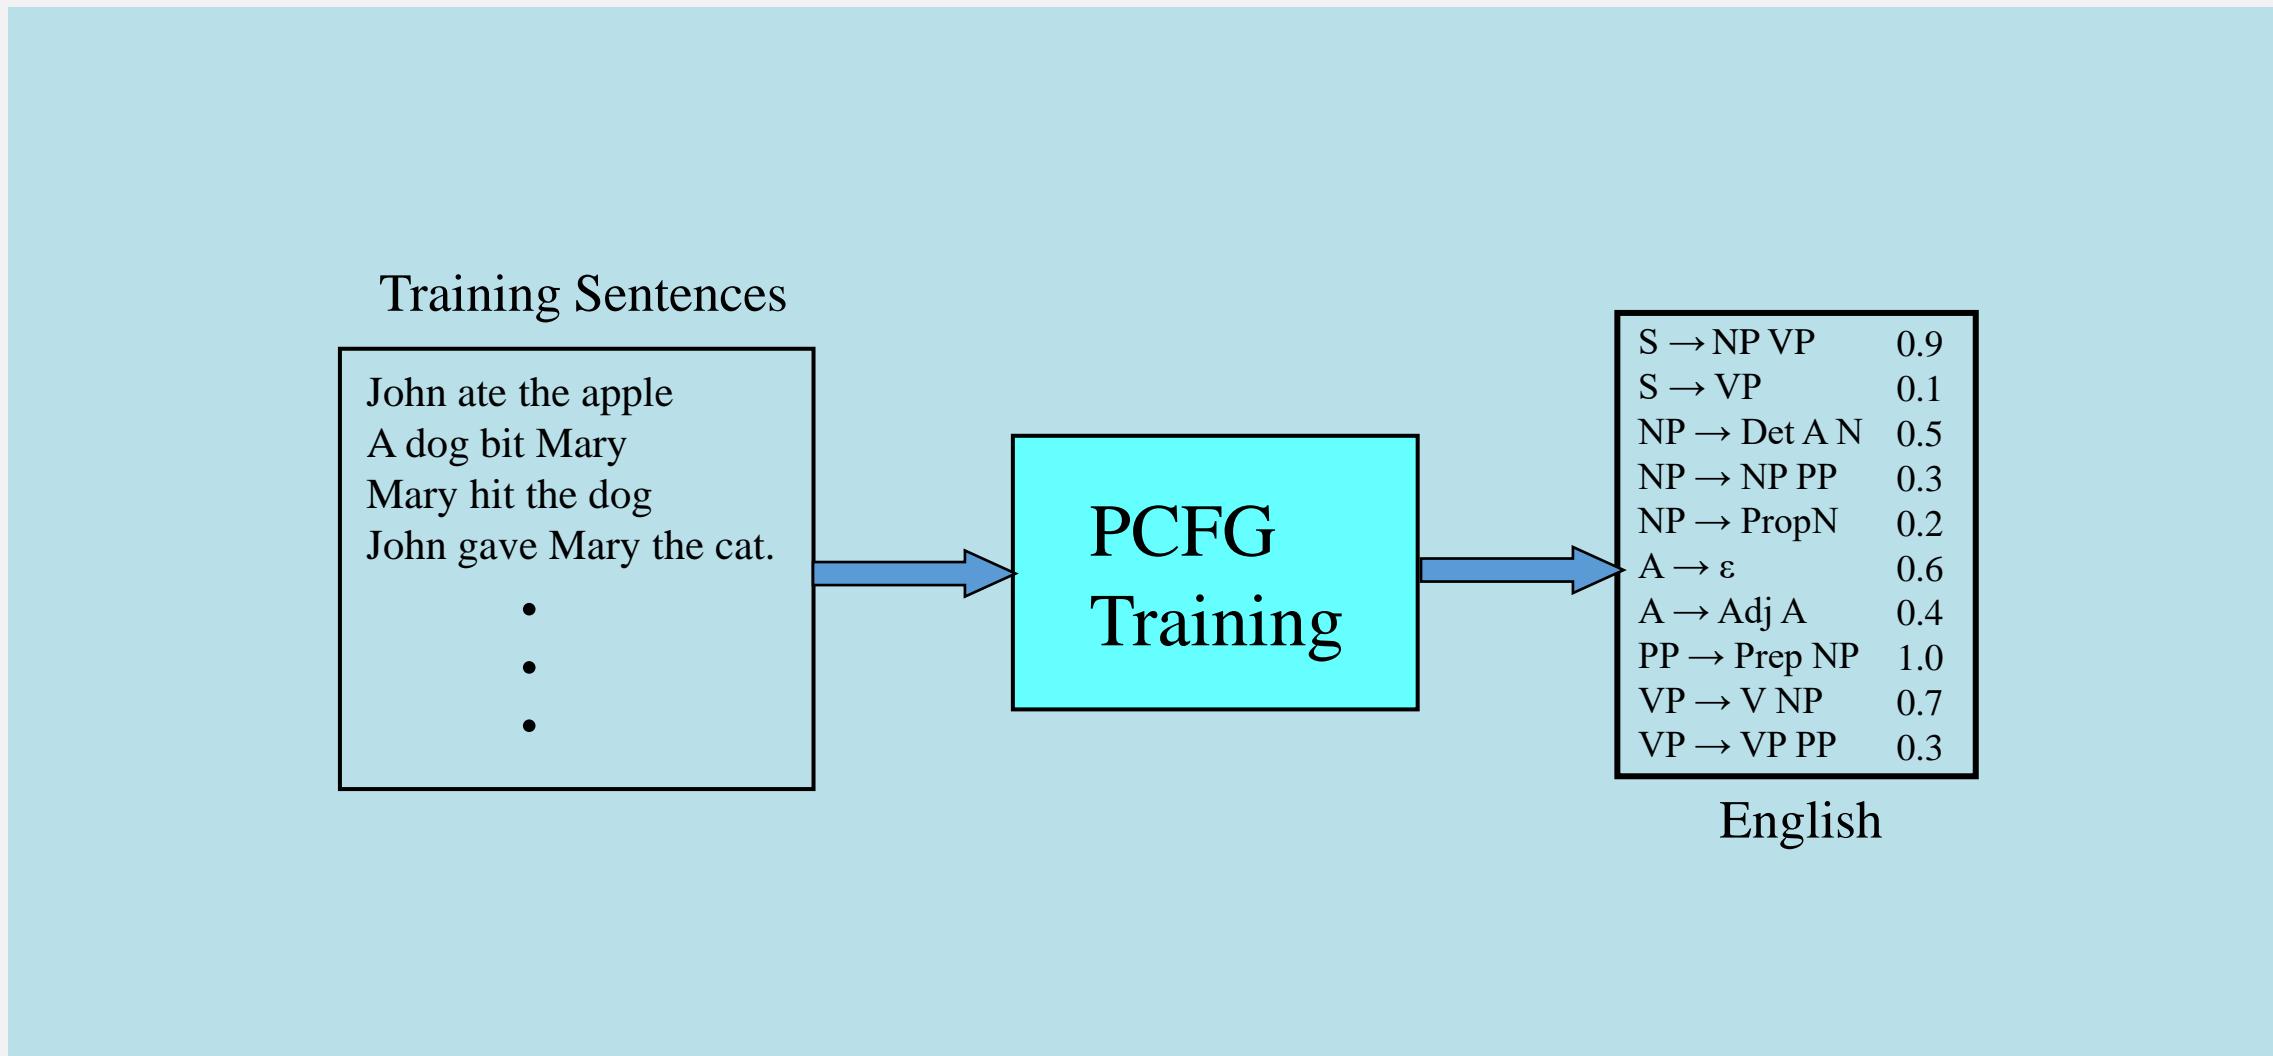
Adapted from Russell, S., & Norvig, P. (2016)

9.2 PCFG: Maximum Likelihood Training

- Given a set of sentences, induce a grammar that maximizes the probability that this data was generated from this grammar.
- Assume the number of non-terminals in the grammar is specified.
- Only need to have an unannotated set of sequences generated from the model. Does not need correct parse trees for these sentences. In this sense, it is unsupervised.

Adapted from Russell, S., & Norvig, P. (2016)

9.2 PCFG: Maximum Likelihood Training



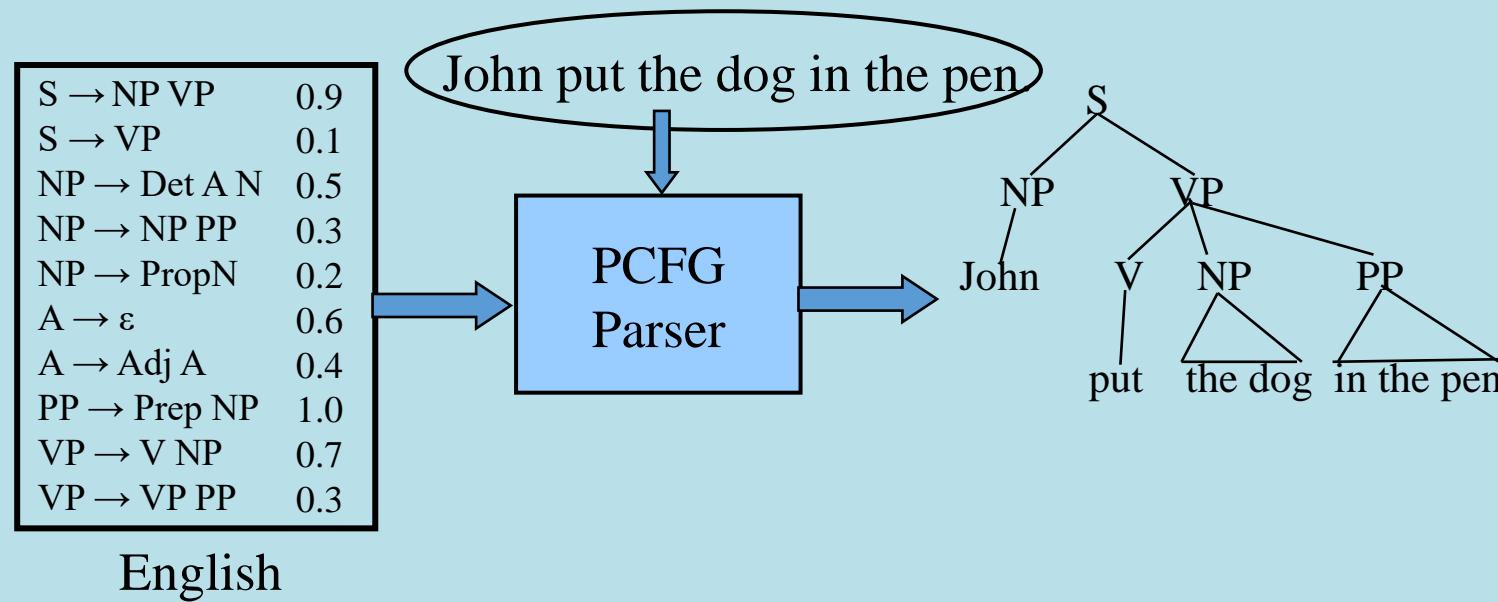
Adapted from Russell, S., & Norvig, P. (2016)

9.2 Vanilla PCFG Limitations

- Since probabilities of productions do not rely on specific words or concepts, only general structural disambiguation is possible (e.g. prefer to attach PPs to Nominals).
- Consequently, vanilla PCFGs cannot resolve syntactic ambiguities that require semantics to resolve, e.g. ate with fork vs. meatballs.
- In order to work well, PCFGs must be lexicalized, i.e. productions must be specialized to specific words by including their head-word in their LHS non-terminals (e.g. VP-ate).

9.2 Example of Importance of Lexicalization

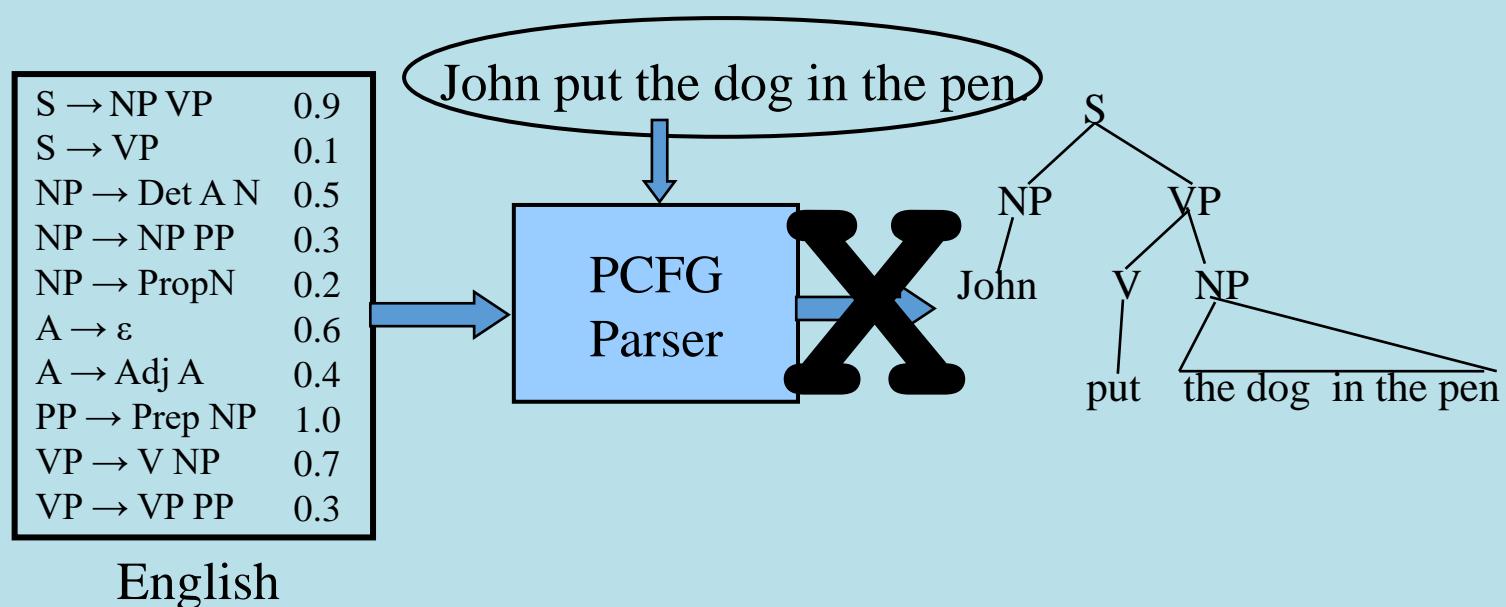
- A general preference for attaching PPs to NPs rather than VPs can be learned by a vanilla PCFG.
- But the desired preference can depend on specific words.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Example of Importance of Lexicalization

- A general preference for attaching PPs to NPs rather than VPs can be learned by a vanilla PCFG.
- But the desired preference can depend on specific words.



Adapted from Russell, S., & Norvig, P. (2016)

9.2 Treebanks

- English Penn Treebank: Standard corpus for testing syntactic parsing consists of 1.2 M words of text from the Wall Street Journal (WSJ).
- Typical to train on about 40,000 parsed sentences and test on an additional standard disjoint test set of 2,416 sentences.
- Chinese Penn Treebank: 100K words from the Xinhua news service.
- Other corpora existing in many languages, see the Wikipedia article “Treebank”

Adapted from Russell, S., & Norvig, P. (2016)

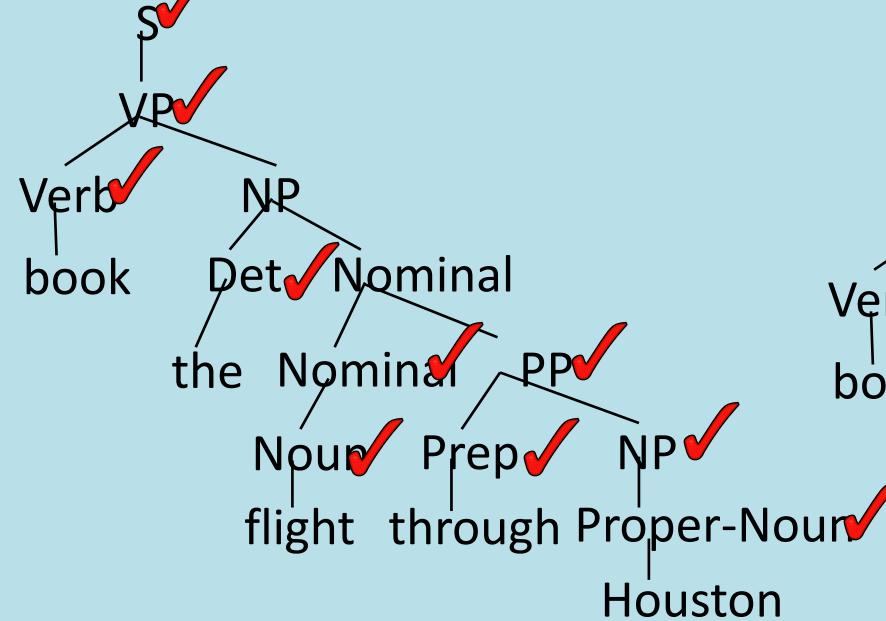
9.2 Parsing Evaluation Metrics

- PARSEVAL metrics measure the fraction of the constituents that match between the computed and human parse trees. If P is the system's parse tree and T is the human parse tree (the “gold standard”):
 - Recall = (# correct constituents in P) / (# constituents in T)
 - Precision = (# correct constituents in P) / (# constituents in P)
- Labeled Precision and labeled recall require getting the non-terminal label on the constituent node correct to count as correct.
- F1 is the harmonic mean of precision and recall.

Adapted from Russell, S., & Norvig, P. (2016)

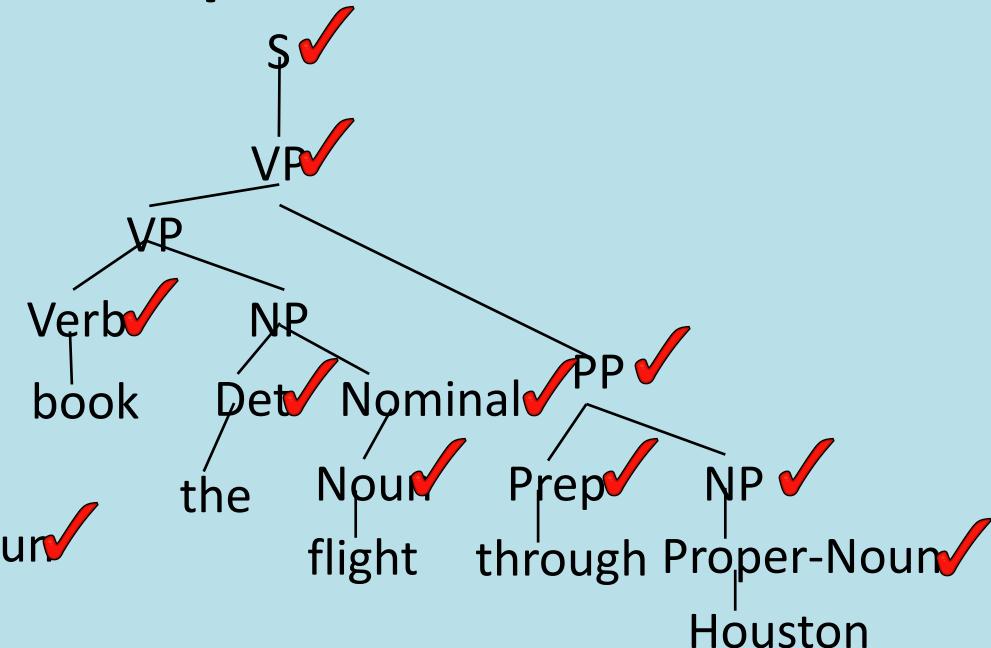
9.2 Computing Evaluation Metrics

Correct Tree T



Constituents: 12

Computed Tree P



Constituents: 12

Correct Constituents: 10

Recall = $10/12 = 83.3\%$ Precision = $10/12 = 83.3\%$

$F_1 = 83.3\%$

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Treebank Results

- Results of current state-of-the-art systems on the English Penn WSJ treebank are slightly greater than 90% labeled precision and recall.

9.2 Word Sense Disambiguation (WSD)

- Words in natural language usually have a fair number of different possible meanings.
 - Ellen has a strong interest in computational linguistics.
 - Ellen pays a large amount of interest on her credit card.
- For many tasks (question answering, translation), the proper sense of each ambiguous word in a sentence must be determined.

9.2 Ambiguity Resolution is Required for Translation

- Syntactic and semantic ambiguities must be properly resolved for correct translation:
 - “John plays the guitar.” → “John toca la guitarra.”
 - “John plays soccer.” → “John juega el fútbol.”
- An apocryphal story is that an early MT system gave the following results when translating from English to Russian and then back to English:
 - “The spirit is willing but the flesh is weak.” → “The liquor is good but the meat is spoiled.”
 - “Out of sight, out of mind.” → “Invisible idiot.”

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Word Sense Disambiguation (WSD) as Text Categorization

- Each sense of an ambiguous word is treated as a category.
 - “play” (verb)
 - play-game
 - play-instrument
 - play-role
 - “pen” (noun)
 - writing-instrument
 - enclosure

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Word Sense Disambiguation (WSD) as Text Categorization

- Treat current sentence (or preceding and current sentence) as a document to be classified.
 - “play”:
 - play-game: “John played soccer in the stadium on Friday.”
 - play-instrument: “John played guitar in the band on Friday.”
 - play-role: “John played Hamlet in the theater on Friday.”
 - “pen”:
 - writing-instrument: “John wrote the letter with a pen in New York.”
 - enclosure: “John put the dog in the pen in New York.”

- Assume part-of-speech (POS), e.g. noun, verb, adjective, for the target word is determined.
- Treat as a classification problem with the appropriate potential senses for the target word given its POS as the categories.
- Encode context using a set of features to be used for disambiguation.
- Train a classifier on labeled data encoded using these features.
- Use the trained classifier to disambiguate future instances of the target word given their contextual features.

9.2 WSD “line” Corpus

- 4,149 examples from newspaper articles containing the word “line.”
- Each instance of “line” labeled with one of 6 senses from WordNet.
- Each example includes a sentence containing “line” and the previous sentence for context.

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Senses of “line”

- Product: “While he wouldn’t estimate the sale price, analysts have estimated that it would exceed \$1 billion. Kraft also told analysts it plans to develop and test a line of refrigerated entrees and desserts, under the Chillary brand name.”
- Formation: “C-LD-R L-V-S V-NNA reads a sign in Caldor’s book department. The 1,000 or so people fighting for a place in line have no trouble filling in the blanks.”
- Text: “Newspaper editor Francis P. Church became famous for a 1897 editorial, addressed to a child, that included the line “Yes, Virginia, there is a Santa Clause.”
- Cord: “It is known as an aggressive, tenacious litigator. Richard D. Parsons, a partner at Patterson, Belknap, Webb and Tyler, likes the experience of opposing Sullivan & Cromwell to “having a thousand-pound tuna on the line.”
- Division: “Today, it is more vital than ever. In 1983, the act was entrenched in a new constitution, which established a tricameral parliament along racial lines, whith separate chambers for whites, coloreds and Asians but none for blacks.”
- Phone: “On the tape recording of Mrs. Guba's call to the 911 emergency line, played at the trial, the baby sitter is heard begging for an ambulance.”

Adapted from Rusell, S., & Norvig, P. (2016)

9.2 Experimental Data for WSD of “line”

- Sample equal number of examples of each sense to construct a corpus of 2,094.
- Represent as simple binary vectors of word occurrences in 2 sentence context.
 - Stop words eliminated
 - Stemmed to eliminate morphological variation
- Final examples represented with 2,859 binary word features.

Adapted from Russell, S., & Norvig, P. (2016)

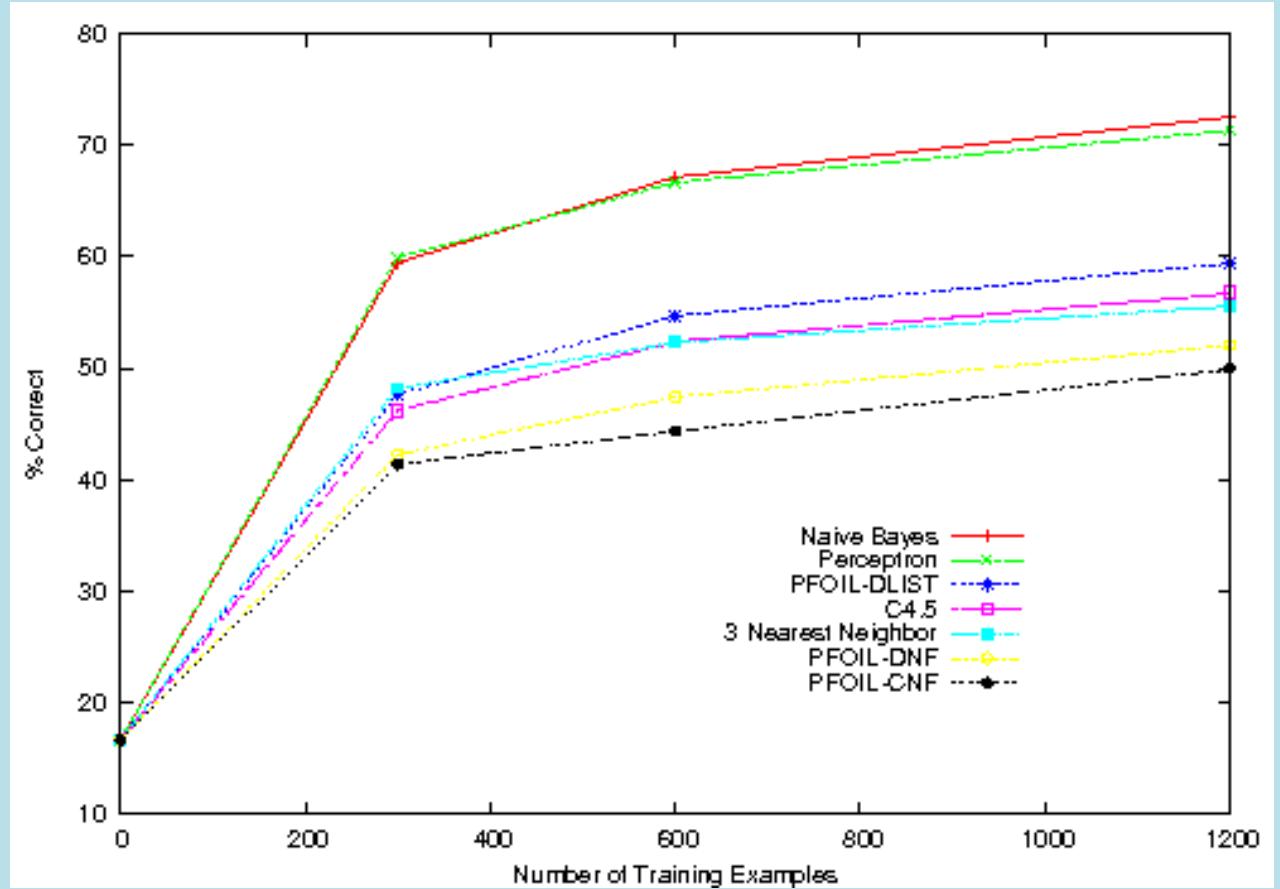
9.2 Machine Learning Algorithms

- Naïve Bayes
 - Binary features
- K Nearest Neighbor
 - Simple instance-based algorithm with k=3 and Hamming distance
- Perceptron
 - Simple neural-network algorithm.
- C4.5
 - State of the art decision-tree induction algorithm
- PFOIL-DNF
 - Simple logical rule learner for Disjunctive Normal Form
- PFOIL-CNF
 - Simple logical rule learner for Conjunctive Normal Form
- PFOIL-DLIST
 - Simple logical rule learner for decision-list of conjunctive rules

Adapted from Russell, S., & Norvig, P. (2016)

9.2 Learning Curves for WSD of “line”

- Naïve Bayes and Perceptron give the best results.
- Both use a weighted linear combination of evidence from many features.
- Symbolic systems that try to find a small set of relevant features tend to overfit the training data and are not as accurate.
- Nearest neighbor method that weights all features equally is also not as accurate.
- Of symbolic systems, decision lists work the best.



Adapted from Russell, S., & Norvig, P. (2016)

Your turn!

Task

Please explain

- What is PCFG
- three tasks of PCFG of your choice

Outline

9 Language and Image Processing

9.1 Applied Natural Language Processing

9.2 Speech Processing and Communication

9.3 Perception and Image Processing

9.4 Generative Modelling

Lectorial 7: Building NLP and Machine Learning Pipelines for Webservices

► What we will learn:

- We will discuss how to make use of the copious knowledge that is expressed in natural language
- And how to process and prepare language data for machines to build state-of-the-art machine learning pipelines for language processing
- We learn new type of models, so called generative models, which allow to generate new instances based on the knowledge of your model



Image source: [↗ Pixabay](#) (2021) / [↗ CC0](#)

► Duration:

- 135 min + 90 (Lectorial)

► Relevant for Exam:

- 9.1, 9.4

Outline

9

Language and Image Processing

9.1 Applied Natural Language Processing

9.2 Speech Processing and Communication

9.3 Perception and Image Processing

9.4 Generative Modelling

Lectorial 7: Building NLP and Machine Learning Pipelines

► **What we will learn:**

- We will discuss how to make use of the copious knowledge that is expressed in natural language
- And how to process and prepare language data for machines to build state-of-the-art machine learning pipelines for language processing
- We learn new type of models, so called generative models, which allow to generate new instances based on the knowledge of your model



Image source: [↗ Pixabay](#) (2021) / [↗ CC0](#)

► **Duration:**

- 135 min + 90 (Lectorial)

► **Relevant for Exam:**

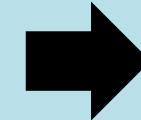
- 9.1, 9.4

9.4 Let us Look behind to Lecture 1



AI model autonomous driving

- **Task:** Driving on public, highway using vision sensors
- **Measure:** Distance before error
- **Training:** sequence of images/video data of human drivers



Machine Learning

A computer program is said to learn from experience ‘E’, with respect to some class of tasks ‘T’ and performance measure ‘P’ if its performance at tasks in ‘T’ as measured by ‘P’ improves with experience ‘E’. (Mitchel, 2011)



**„What I cannot create
I do not understand“**

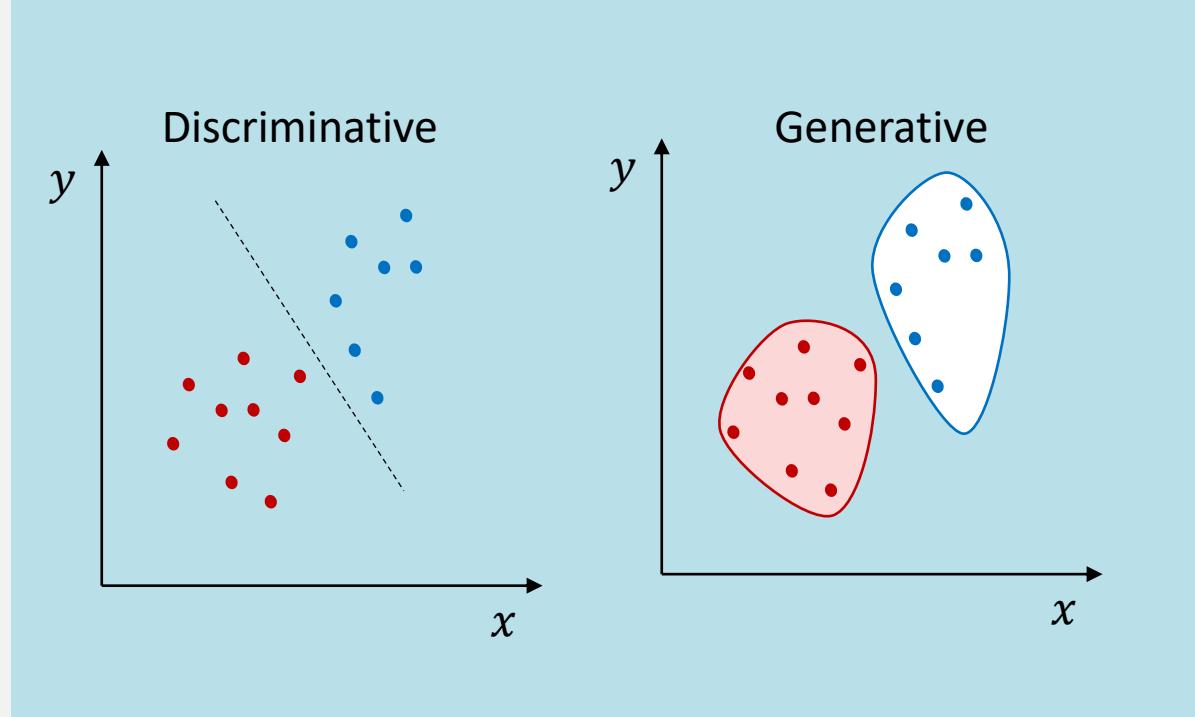
Richard Feynman

Image source: ↗ [Richard Feynman, 1984](#) (1984) by Tamiko Thiel ↗ [CC BY-SA 3.0](#)

9.4 Problem: High-Performance vs. Understanding



Image source: [↗ DB2018AL00555 \(VW\)](#) | free for editorial purposes



Example: Pattern recognition vs. generation in autonomous driving

- Sometimes we have to understand/model the true distribution of our dependent variable
- We want to generate new data instances

Adapted from Ng AY & Jordan MI (2002)

9.4 Formalization: Discriminative Models vs. Generative Models

- We defined machine learning as:

$$f: X \rightarrow Y$$

- To predict the class Y from the training example X in machine learning problem (e.g. classification with decision tree or perceptron), we have to solve

$$f(X) = \arg \max_y p(Y | X)$$

- We try to model this conditional probability by modelling a decision boundy between the classes

9.4 Formalization: Discriminative Models vs. Generative Models

- If we use Bayes' rule we can replace $p(Y | X)$, and get

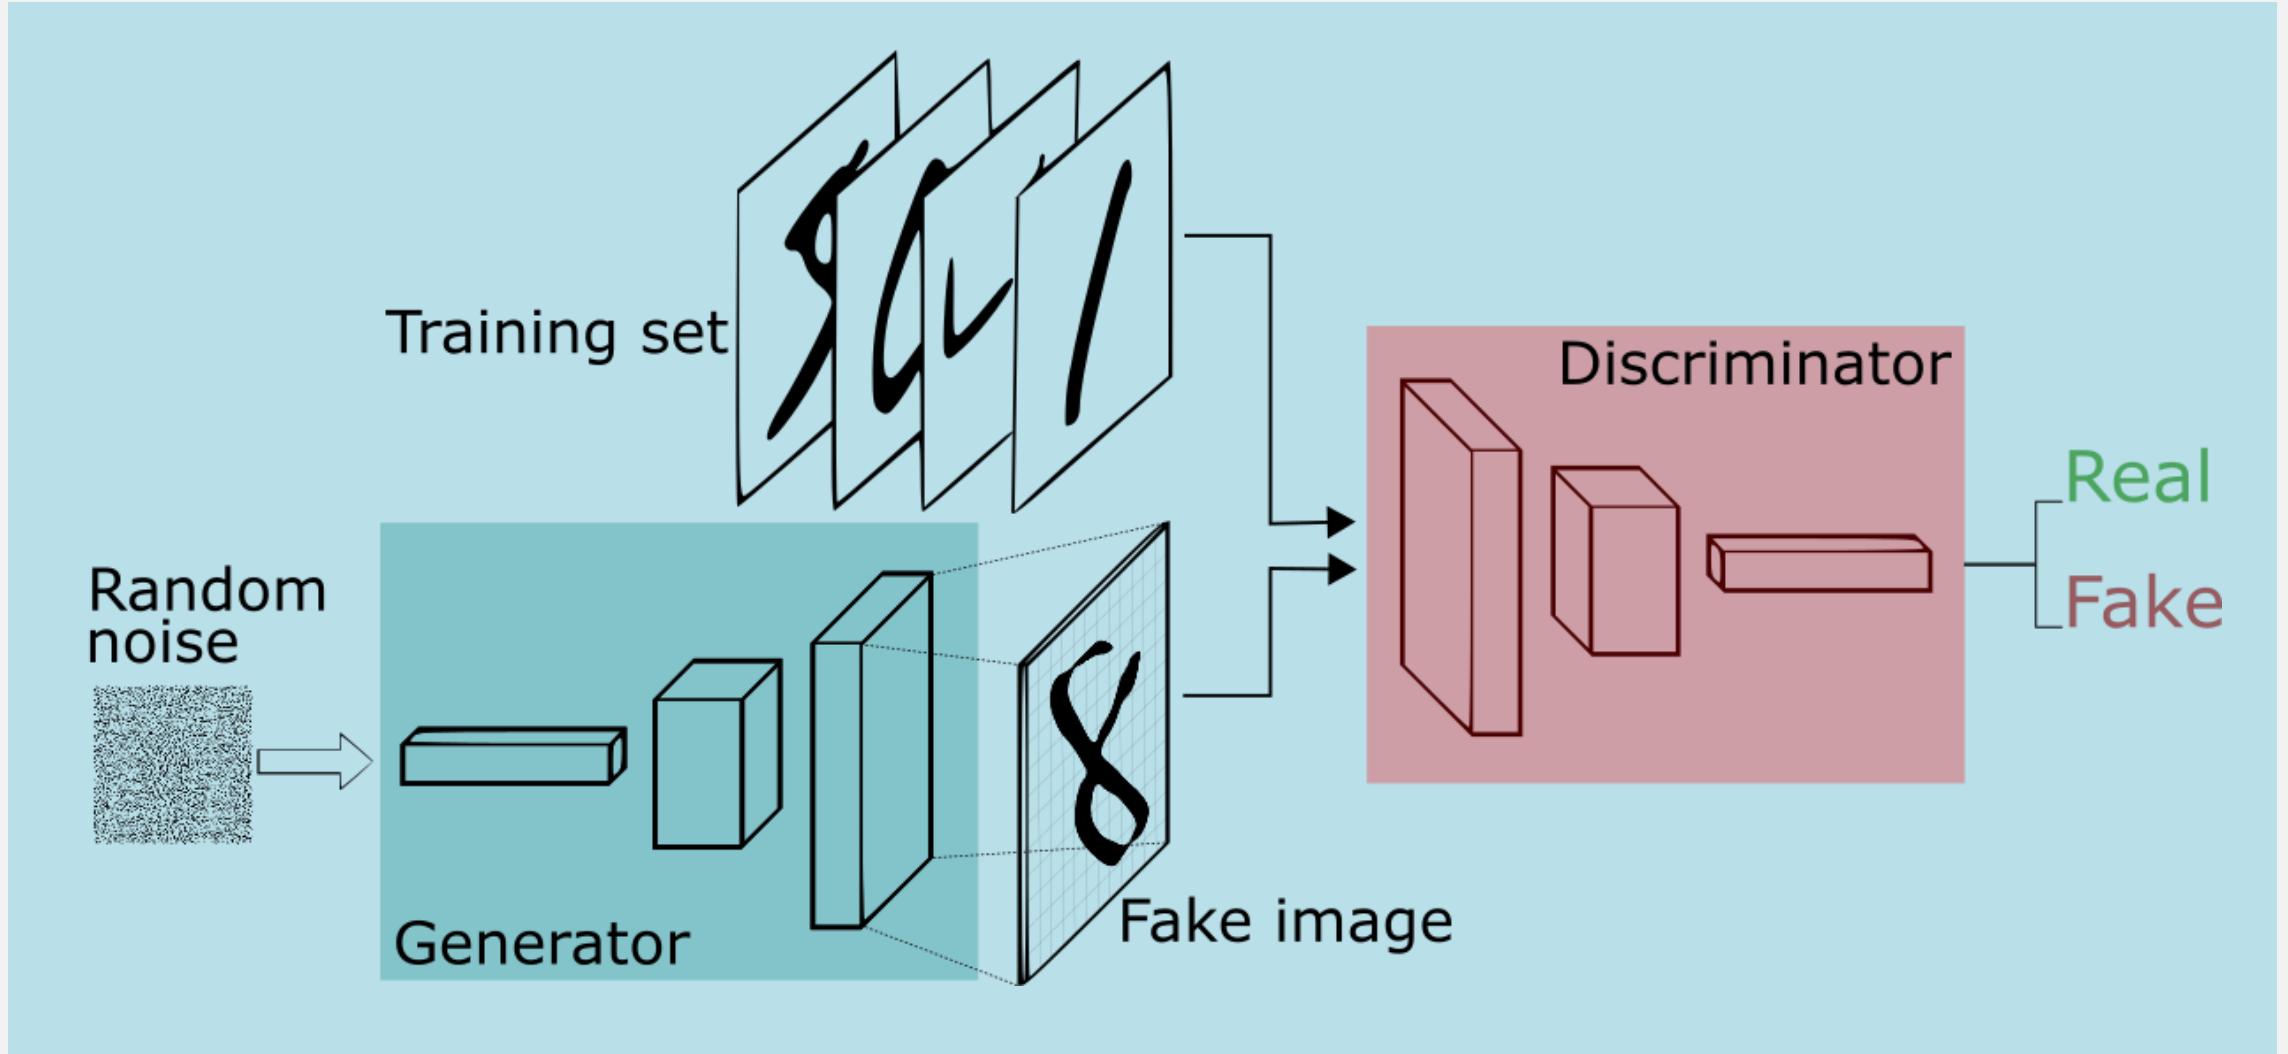
$$f(X) = \arg \max_y \frac{p(X | Y) \cdot p(Y)}{p(X)} \approx \arg \max_y p(X | Y) \cdot p(Y)$$

- With

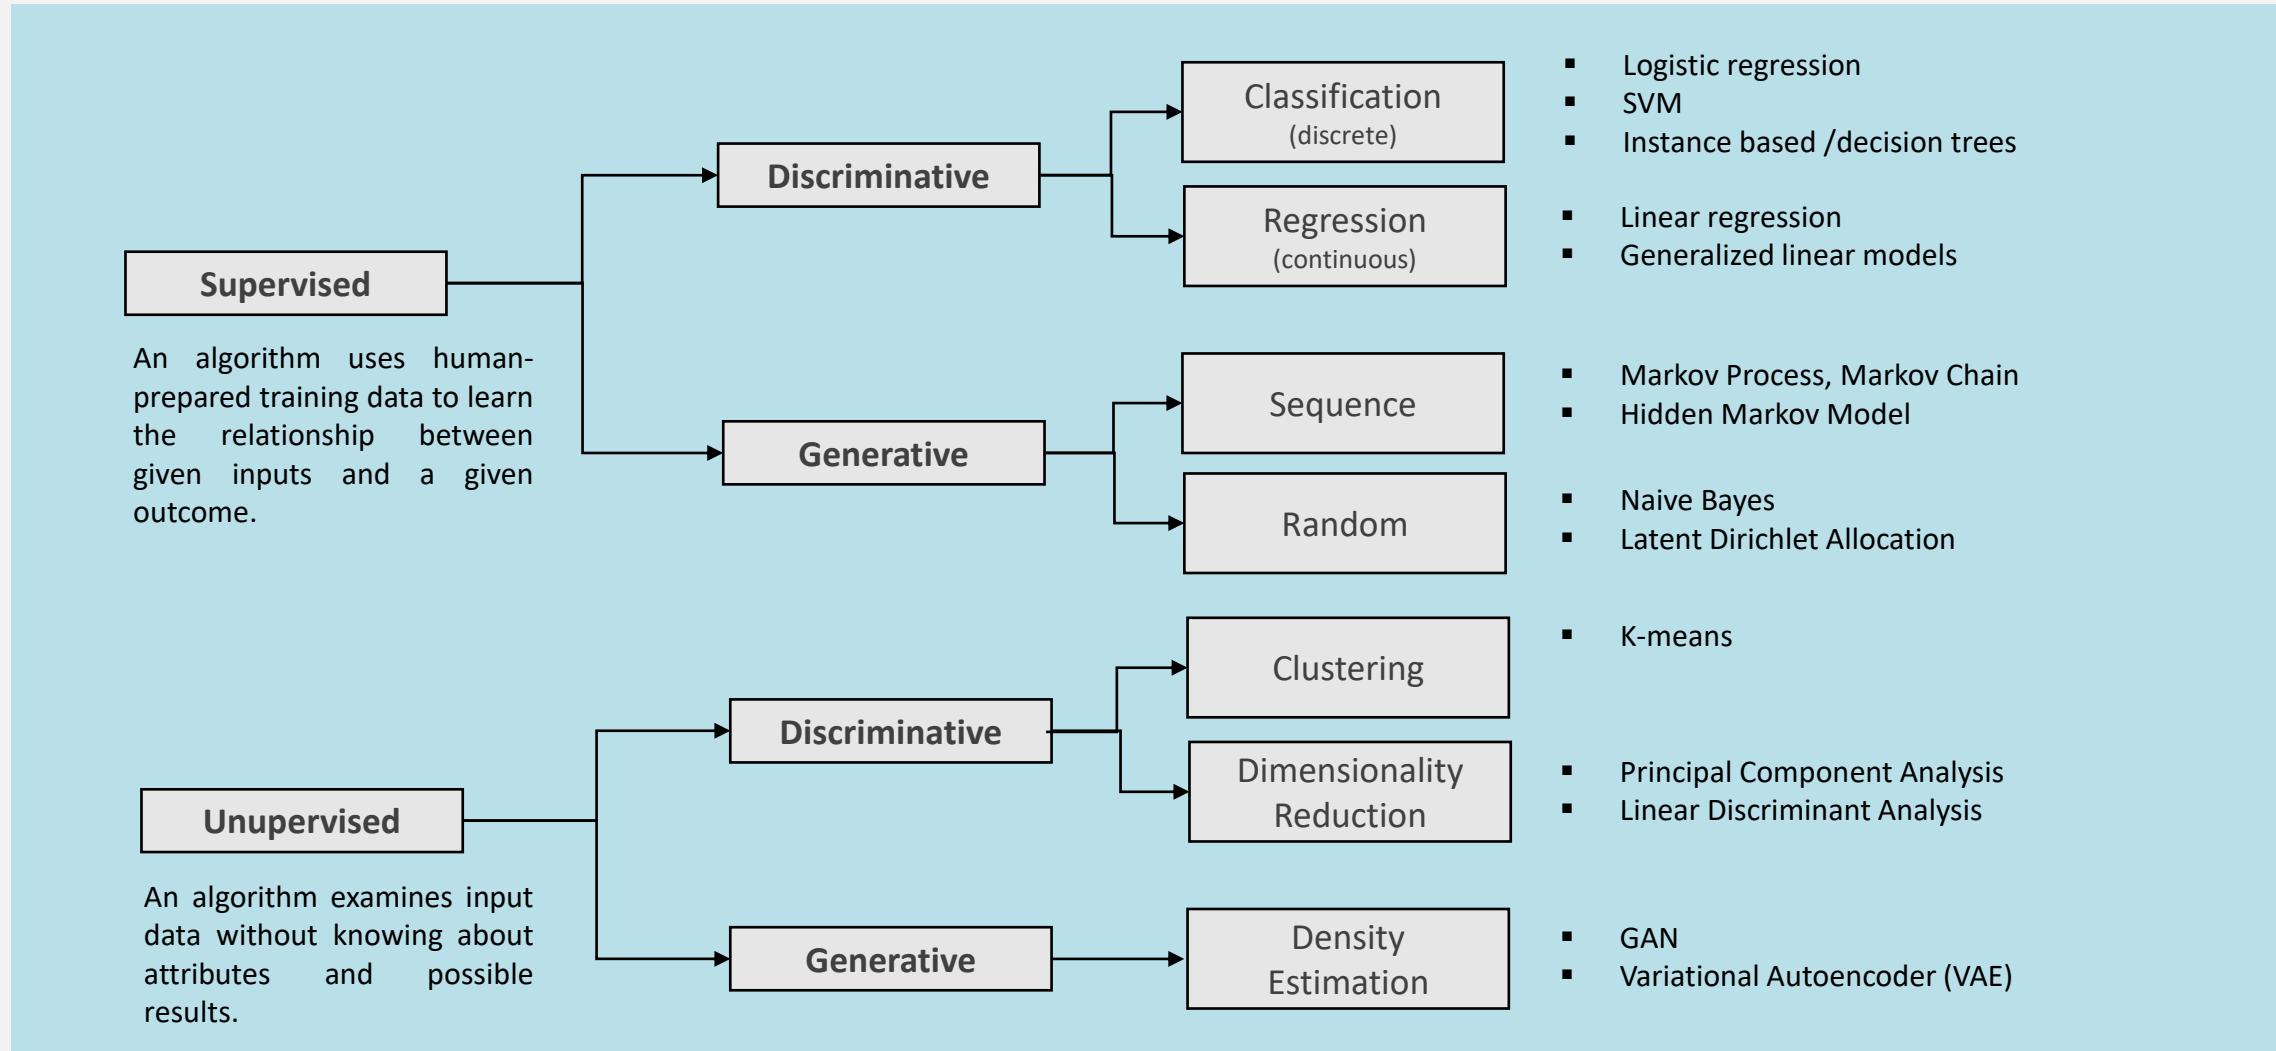
$$p(X | Y) \cdot p(Y) = p(X | Y)$$

- Hence, our problem describes now the joint probability distribution $p(X | Y)$, which models the actual distribution of each class
- These models are called generative, because they allow to calculate/generate the respective X for each Y

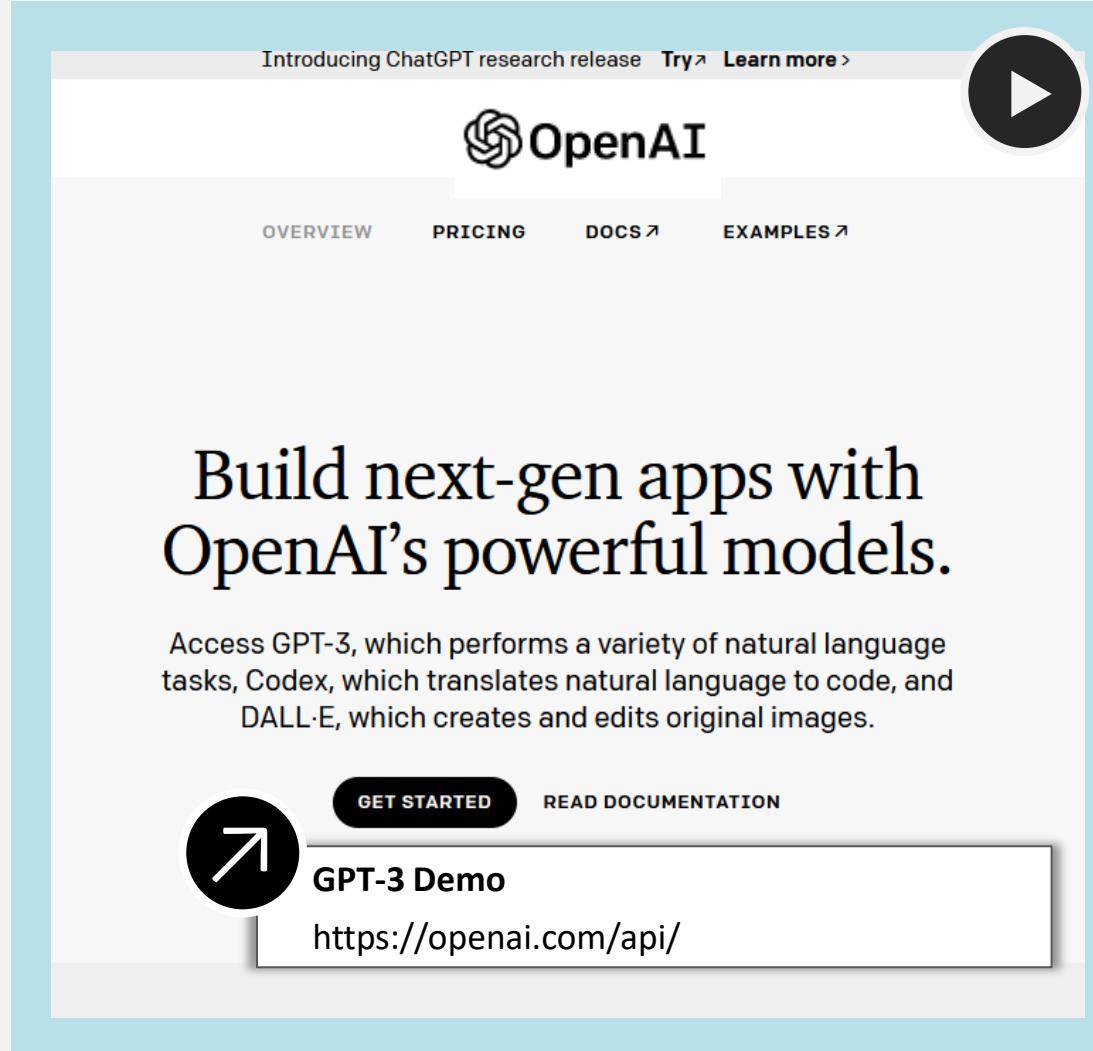
9.4 Generative AI



9.4 Taxonomy: Discriminative vs. Generative in Machine Learning



9.4 Motivation: Generative Pre-trained Transformer (GPT)



The screenshot shows the official OpenAI website. At the top, there's a navigation bar with links for "Introducing ChatGPT research release", "Try", "Learn more", the OpenAI logo, and menu items "OVERVIEW", "PRICING", "DOCS", and "EXAMPLES". A large play button icon is positioned in the top right corner. The main headline reads "Build next-gen apps with OpenAI's powerful models." Below it, text describes three models: GPT-3 (natural language tasks), Codex (code generation), and DALL-E (image creation). At the bottom, there's a "GET STARTED" button with an upward arrow icon, a "READ DOCUMENTATION" link, a "GPT-3 Demo" section with a "GPT-3 Demo" button and a URL "https://openai.com/api/", and a footer with the OpenAI logo and the text "OpenAI © 2023".

- GPT-3 is a language processing model developed by the American non-profit organization OpenAI.
- It uses deep learning (see chapter 7) to create, summarize, simplify or translate texts.
- The OpenAI API can be applied to virtually any task that involves understanding or generating natural language or even code

The screenshot shows a news article from the German magazine **stern**. The header includes the **stern** logo, navigation links for **Panorama**, **Politik**, **Kultur**, **Lifestyle**, **Digital**, **Wirtschaft**, **Sport**, **Gesundheit**, **Genuss**, **Reise**, and **Familie**, and a search bar. The main headline reads: **Künstliche Intelligenz: Diese Seite zeigt Menschen, die es gar nicht gibt – und das ist ziemlich gruselig**. Below the headline is a subtext: "Achtung: Was ihr seht, ist nicht echt. Mit künstlicher Intelligenz können Maschinen nämlich unendlich viele Fake-Gesichter erstellen. Die Ergebnisse der beeindruckenden und zugleich gruseligen Technologie zeigen die Entwickler jetzt im Netz." At the bottom left, there is a watermark for **ThisPersonDoesNotExist.com** with a circular arrow icon.

ABO STERN PLUS CRIME PLUS

VIDEO SPIELE GUTSCHEINE

Panorama Politik Kultur Lifestyle Digital Wirtschaft Sport Gesundheit Genuss Reise Familie

Gutscheine

Home > Neon > Wilde Welt > thispersondoesnotexist.com zeigt Menschen, die es gar nicht gibt

Künstliche Intelligenz: Diese Seite zeigt Menschen, die es gar nicht gibt – und das ist ziemlich gruselig

Achtung: Was ihr seht, ist nicht echt. Mit künstlicher Intelligenz können Maschinen nämlich unendlich viele Fake-Gesichter erstellen. Die Ergebnisse der beeindruckenden und zugleich gruseligen Technologie zeigen die Entwickler jetzt im Netz.

ThisPersonDoesNotExist.com

Source. Stern (2019)



Bloomberg the Company & Its Products ▾ | Bloomberg Terminal Demo Request | Bloomberg Anywhere Remote Login | Bloomberg Customer Support

Bloomberg Europe Edition ▾ Sign In Subscribe

• Live Now Markets Economics Industries **Technology** Politics Wealth Pursuits Opinion Businessweek Equality Green CityLab Crypto More

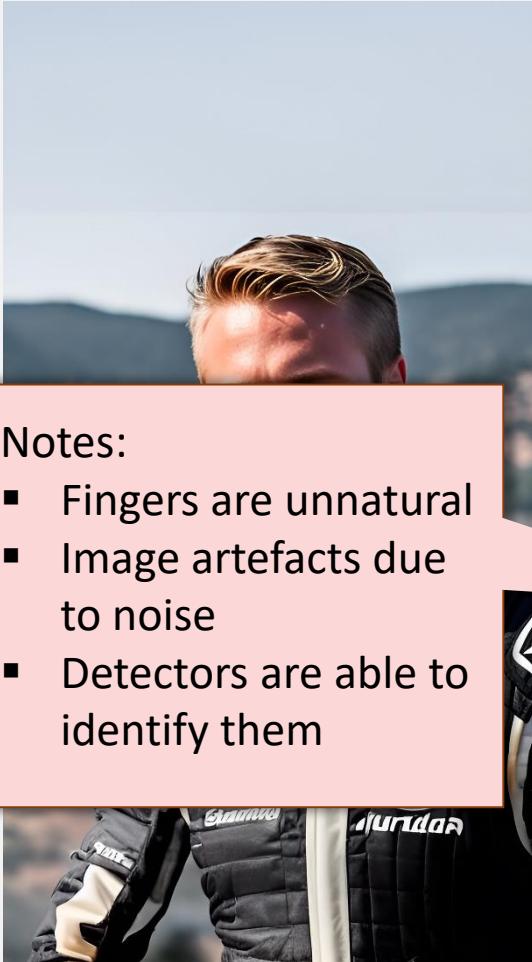
Newsletter

AI Deep Fake of the Pope's Puffy Coat Shows the Power of the Human Mind



Source: Bloomberg (2023)

9.4 Problems and Challenges with Generative Images

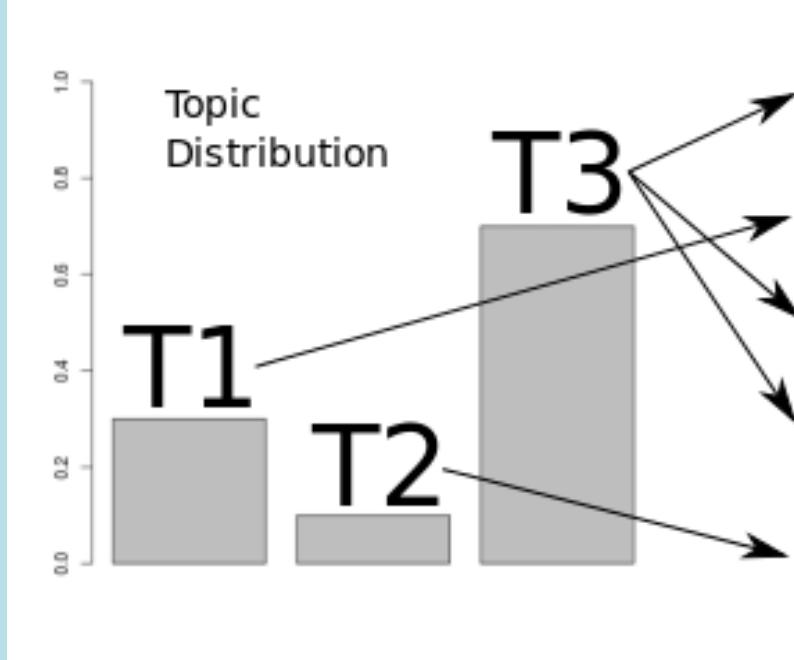


- Fingers are unnatural
- Image artefacts due to noise
- Detectors are able to identify them



9.4 Latent Dirichlet Allocation (LDA)

- Topic model approaches to analyze unstructured data to find hidden topics



Example

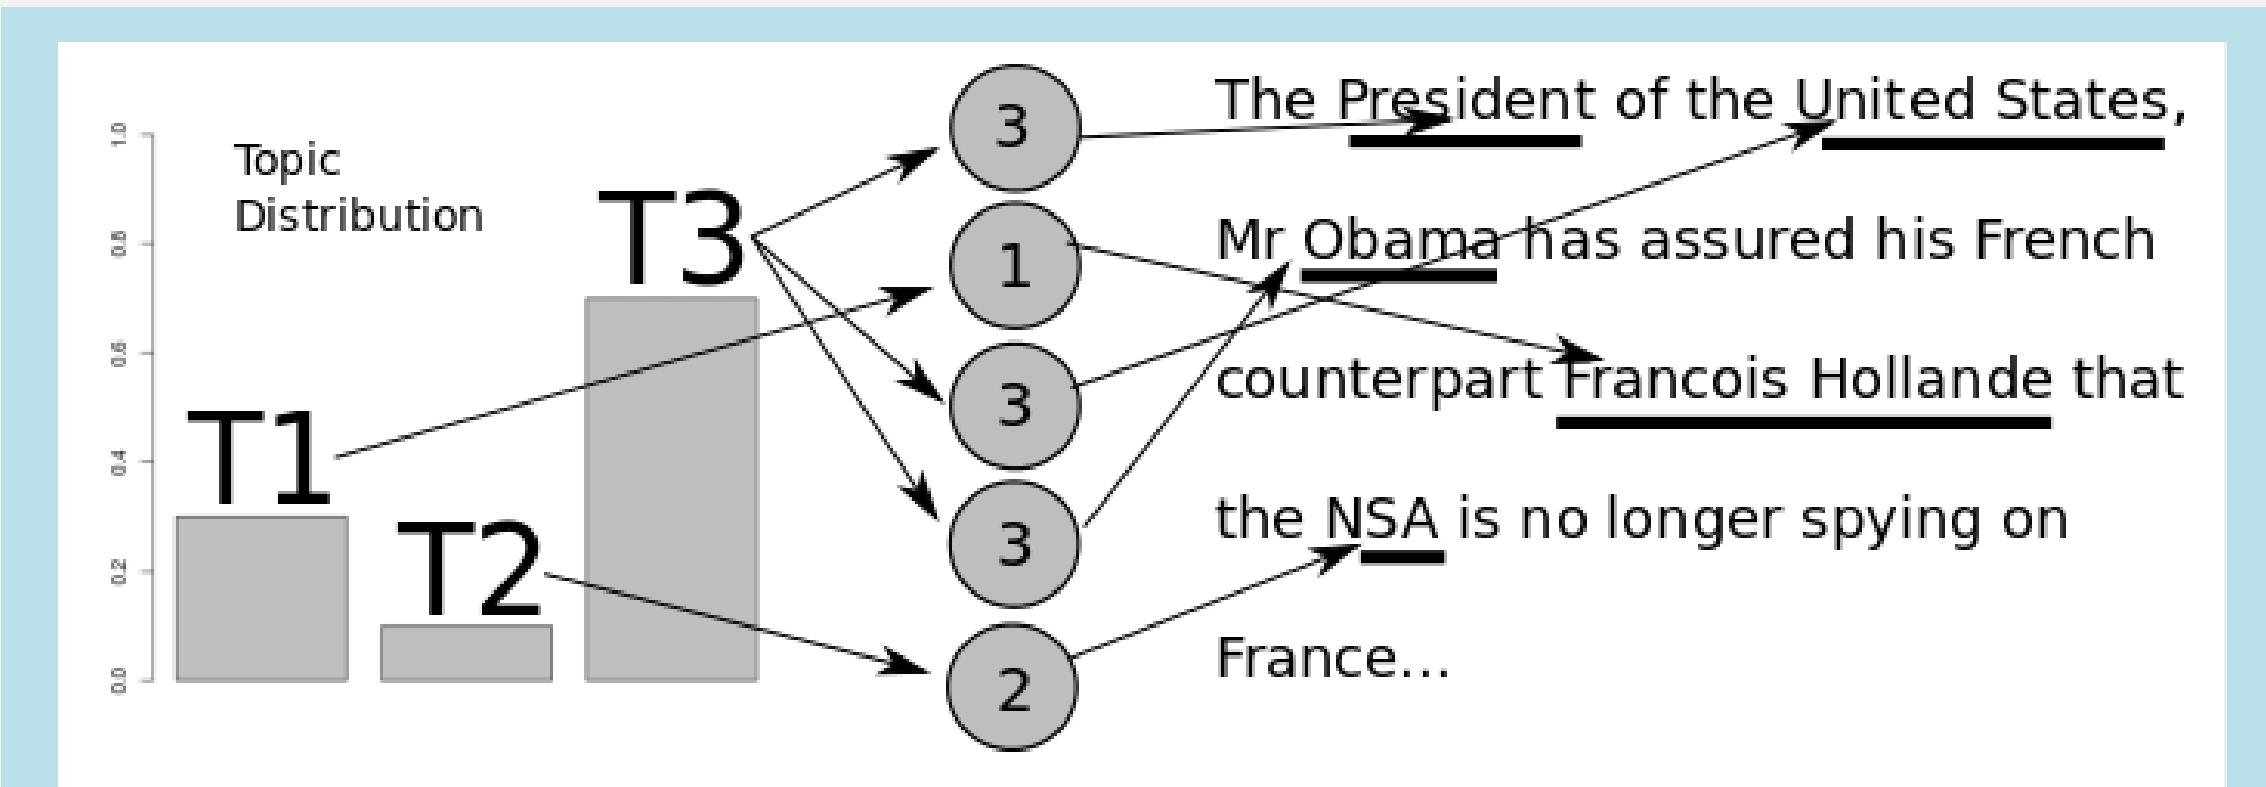
- Automated Summarization, Mediarecommender

- One of the most popular and simplest topic model approaches to analyze unstructured
- Latent Dirichlet allocation is not the best but today by far the most popular approach.
- The method is based on the work of Deerwester in latent semantic indexing and of Hofmann in probabilistic latent semantic indexing

- + probabilistic model with interpretable topics.
- hard to know when LDA is working, because topics are soft-clusters so there is no objective metric

- Idea of latent Dirichlet allocation is to understand the document creation as a stochastic process (see chapter 8)
- It is assumed that a document w_j is generated by an author who selects words of a vocabulary V with a given probability from different baskets of words where each basket corresponds to one of k topics
- The created texts are the observed variables, while the topics, the distribution of the topics and the assignments per document and word is hidden. With that assumptions, the purpose is to decompose from a document with i words, $w = (w_1; w_2; \dots; w_i)$ the original topic baskets

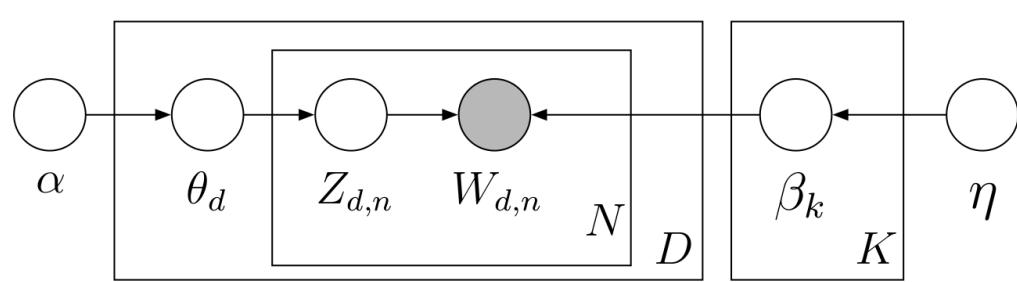
9.4 Topics in LDA



- Each topic contains a individual basket of words and is chosen by his probability (the circles with numbers)
- Each time a topic is chosen, a word ofhis basket is drawn by his probability

Adapted from David M. Blei et al. (2006)

9.4 Latent Dirichlet Allocation



$$p(\vec{\theta}_{1:D}, z_{1:D,1:n}, \vec{\beta}_{1:K} | \omega_{1:D,1:n}, \alpha, \eta) =$$

$$\frac{p(\vec{\theta}_{1:D}, \vec{z}_{1:D}, \vec{\beta}_{1:K} | \omega_{1:D,1:n}, \alpha, \eta)}{\int_{\vec{\beta}_{1:K}} \int_{\vec{\theta}_{1:D}} p(\vec{\theta}_{1:D}, \vec{z}_{1:D}, \vec{\beta}_{1:K} | \omega_{1:D,1:n})}$$

Algorithm: Generative LDA Process

```
for each topic in number of topics k do
    Draw a distribution over words  $\beta_k \sim Dir(\eta)$ 
end for
```

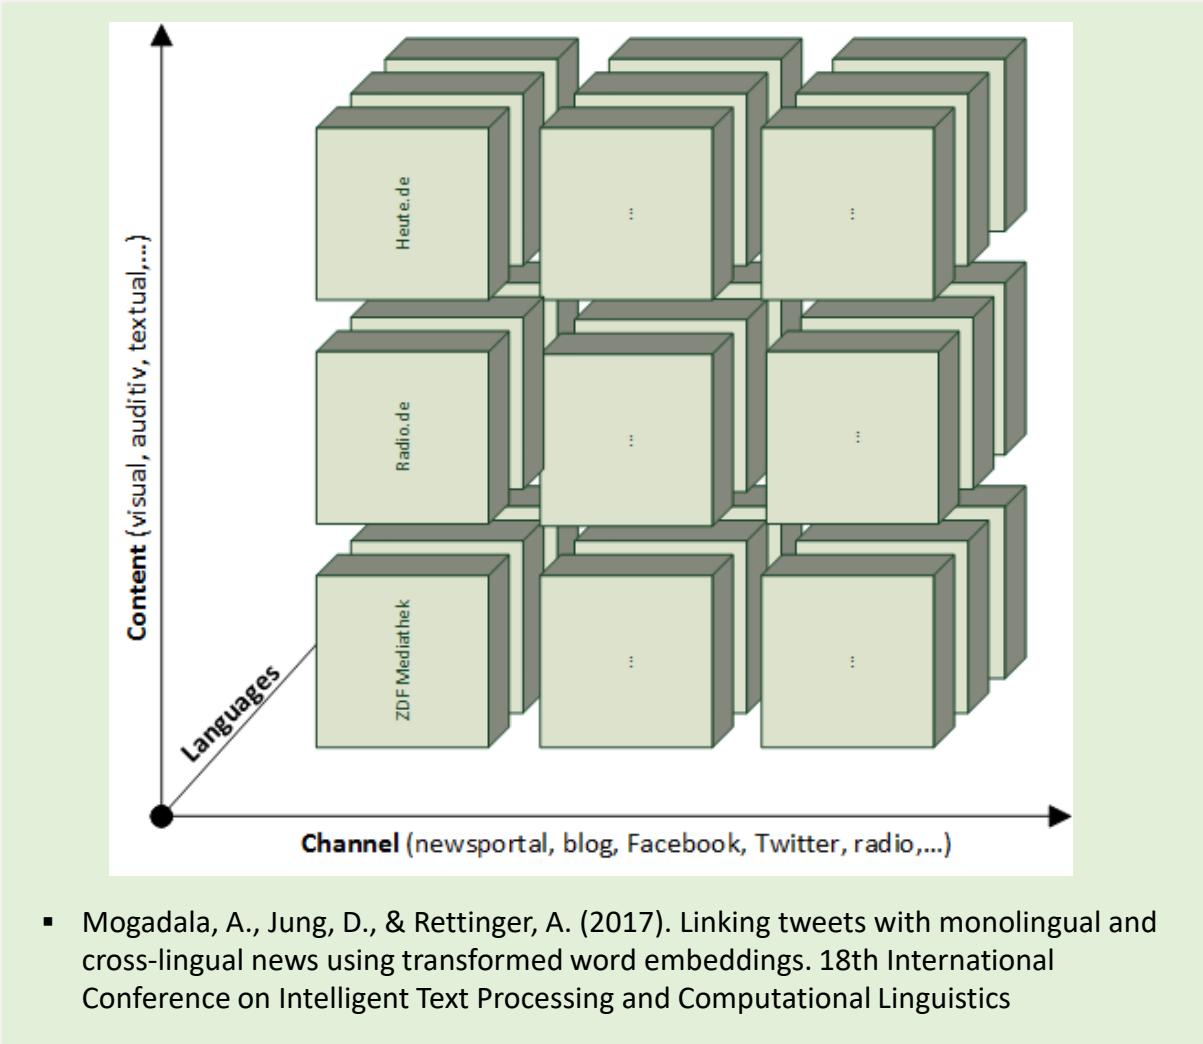
```
for each document  $w_n = d$  in corpus D do
    Draw a vector of topic proportions  $\vec{\theta}_d \sim Dir(\alpha)$ 
    for each word  $w_n^i$  do
        Draw a topic assignment  $Z_{d;n} \sim Mult(\vec{\theta}_d)_{Z_{d,n} \in [1\dots K]}$ 
        Draw a word  $w_{d;n} \sim Mult(\vec{\beta}_{Z_{d,n}})_{w_{d,n} \in [1\dots K]}$ 
    end for
end for
```

Adapted from David M. Blei et al. (2006)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations

Multi-Modal AI is hard

- EU is one of the multilingual markets, with more than 24 official languages
- Today there exists different media items such as email, blog, books, journals, articles in multiple media formats
- Challenge: very heterogeneous media and information. Hence, tasks as information retrieval, searching, monitoring, recommendations, classification, evaluation, translation, summarization or further media analysis pose a major challenge.



Adapted from Mogadala et al. (2017)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations

- Corpus is a collection of n distinct textual media items and is denoted by

$$\mathcal{D} = \{ w^1, w^2, \dots, w^n \}_{n \in \mathbb{N}}$$

- Document of this corpus is a sequence of words:

$$w^j = (w_1^j, w_2^j, \dots, w_i^j)$$

- Aim to find a semantic similarity function $sim()$ or a dissimilarity function $dis()$ to identify relatedness between items. Such a function takes the entities as an input and outputs a score, which indicates the semantic relatedness of the input entity sets.

$$\forall w^j, w^k \in \mathcal{D}: sim(w^j, w^k) = s_{j,k} \in \mathbb{R}_{[0,1]}$$

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations

- Dissimilarity function $dis()$ can be expressed depending on the similarity function $sim()$

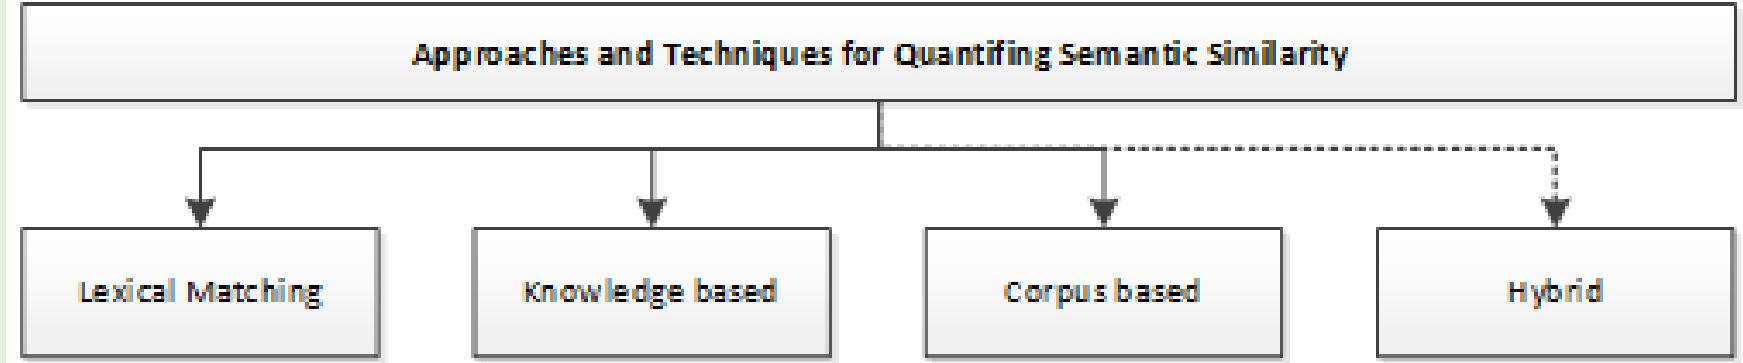
$$dis(w^j, w^k) = 1 - sim(w^j, w^k)$$

range [0,1], 0: no similarity

- **Multi-Modal Recommender:** Find the entity combination of the corpus with the highest similarity

$$\begin{aligned} & \text{Most similar doc pair } (w^j, w^k)^* \\ &= \arg \max_{j,k} sim(w^j, w^k), s.t. w^j, w^k \in \mathcal{D} \end{aligned}$$

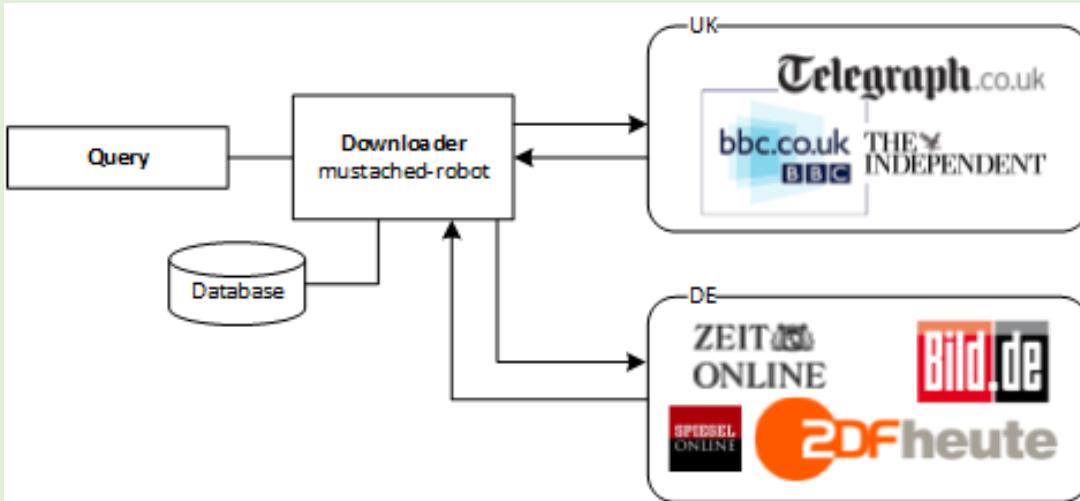
9.4 Business Case: Topic Modelling for Multi-Modal Recommendations



- **Lexical matching methods** include methods which aim to calculate a similarity score by counting lexical units that exits in the input texts.
- **Corpus-based algorithms** calculate by means of large corpora information for a similarity score
- **Knowledge-based** measures calculate a similarity score by means of large networks like Word-Net

Adapted from Mogadala et al. (201/)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations



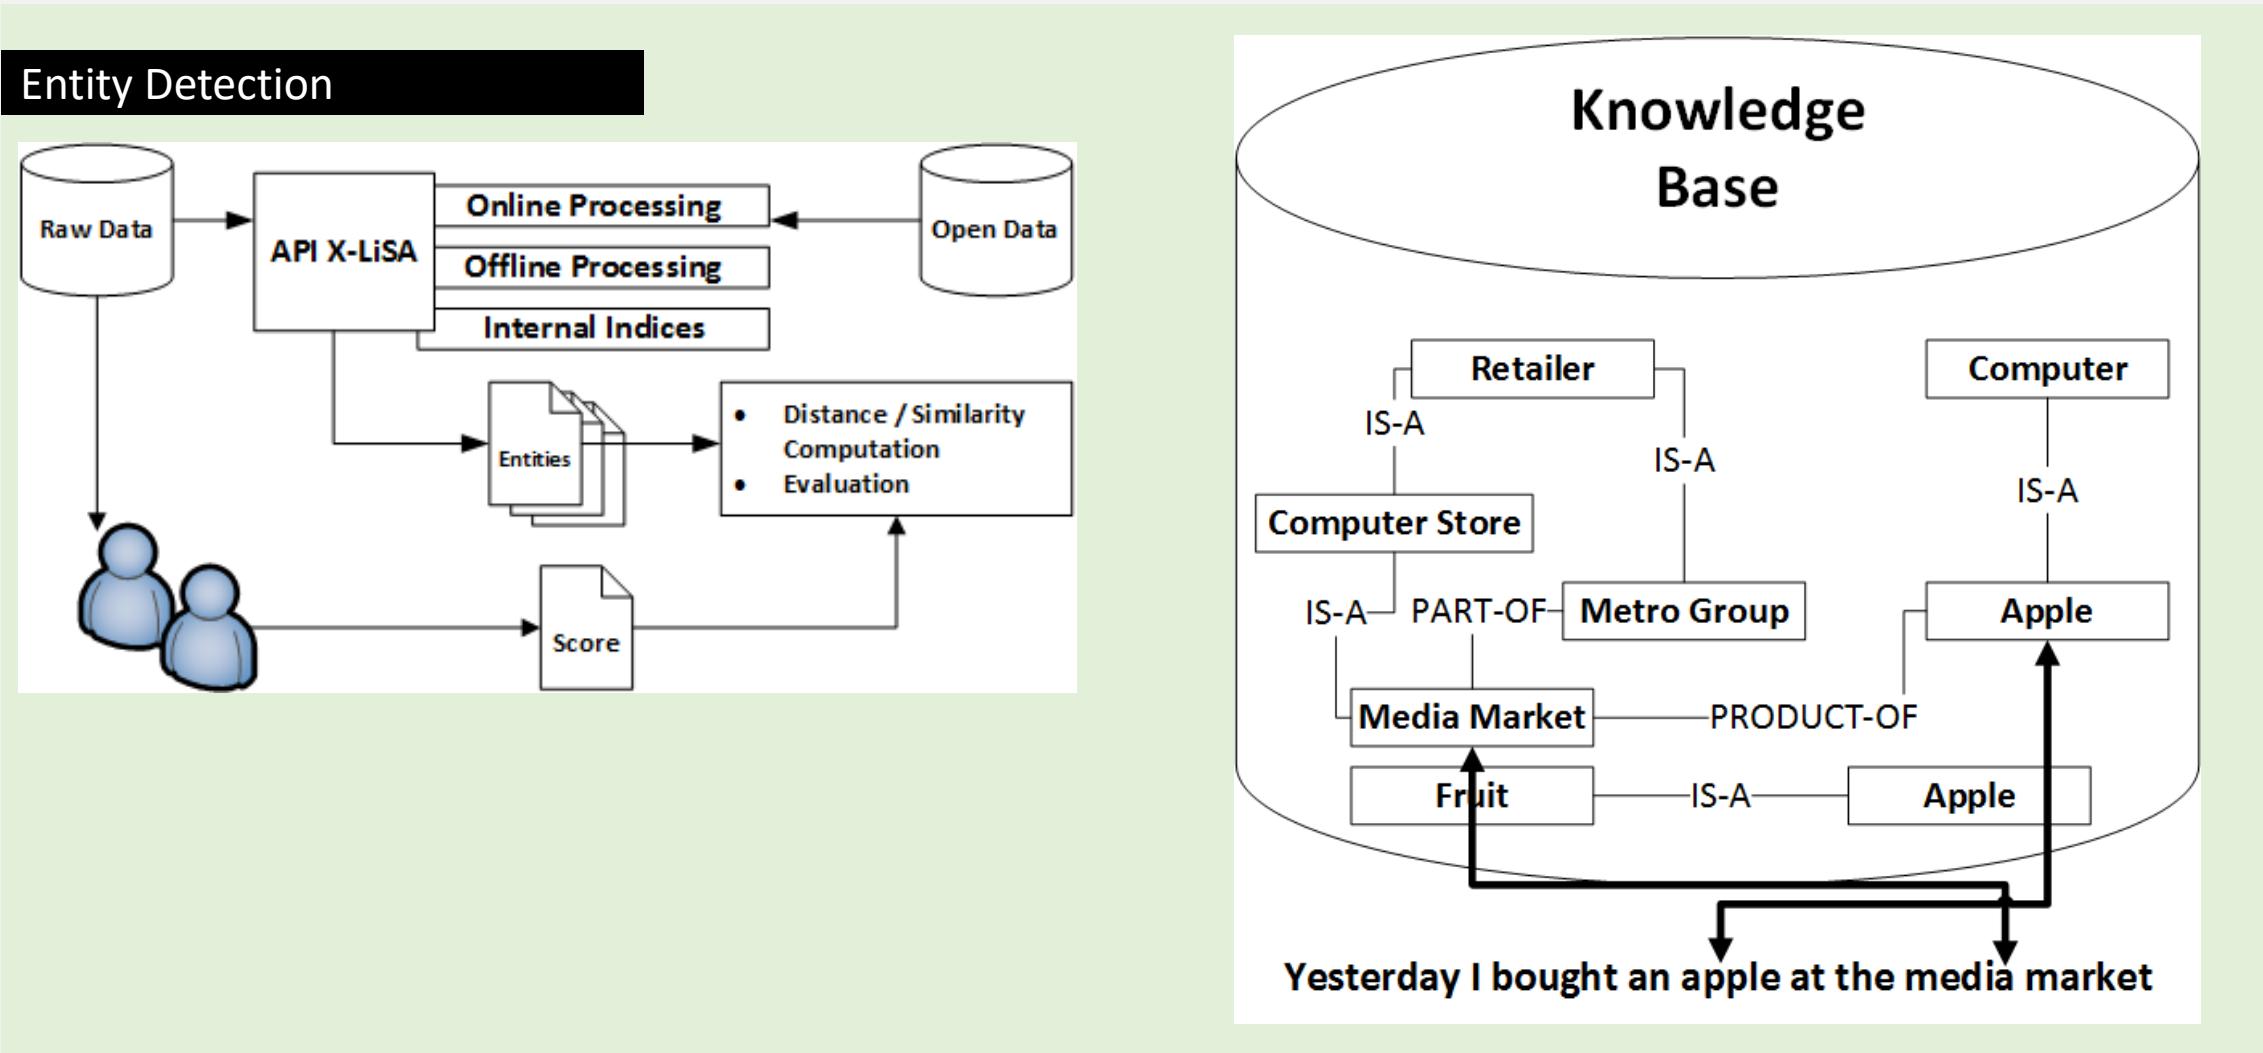
	Dataset 1	Dataset 2
Language of the corpus	Cross-lingual	Cross-lingual
Heterogeneity of the corpus	1 corpus topic	2 corpus topics
Noise of the corpus	Cross-channel	Mono-channel

- To have comparable datasets, all the documents were crawled by two different keywords "Grexit" (the greek exit of the EU) and "4U9525" (airplain crash in france) from the newsportal from January 2015 until May 2015

- To evaluate our model we conducted a user study. Each participant had to rate all combinations of articles for semantic similarity pairwise

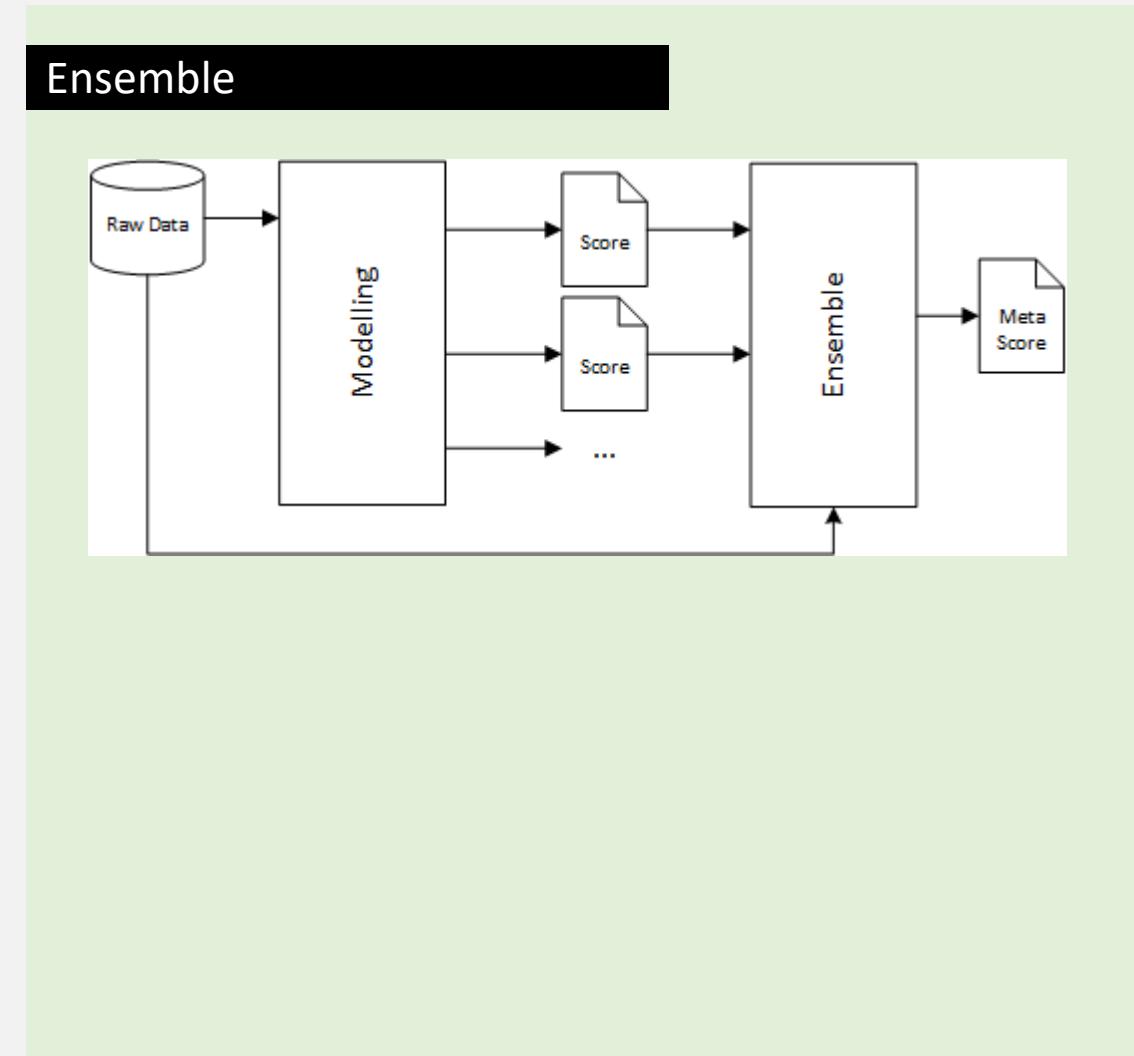
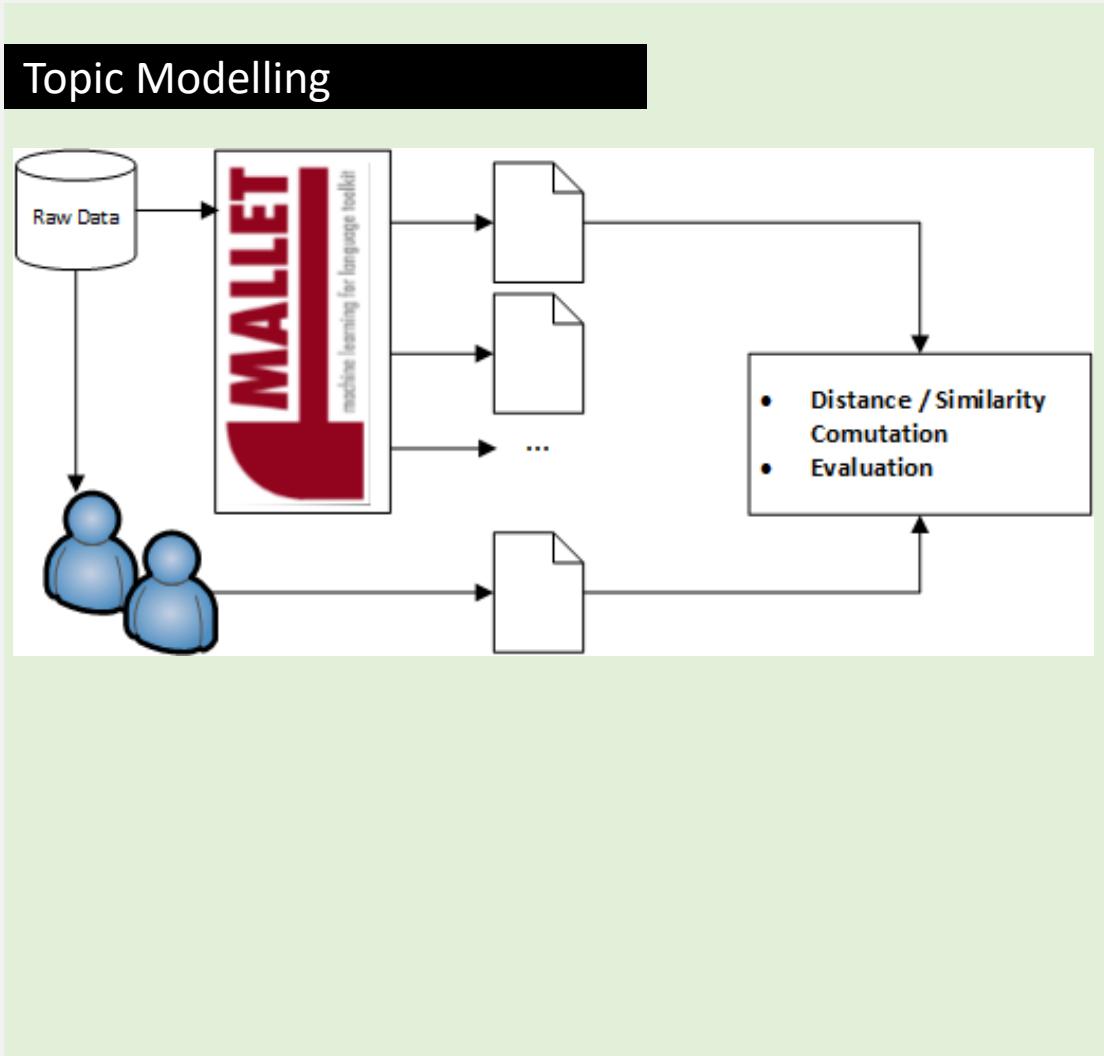
Adapted from Mogadala et al. (201/)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations



Adapted from Mogadala et al. (201/)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations



Adapted from Mogadala et al. (201/)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations

Method and similarity function	Accuracy	
	Dataset 1	Dataset 2
Topic modelling: Euclidean	28.4%	27.72%
Topic modelling: Hellinger	33.64%	37.72%
Entity linking: Subset	33.95%	91.01%
Entity linking: Jaccard	35.18%	94.58%
Ensemble learning algorithm	Dataset 1	Dataset 2
	Ranking learning	66.98%
C4.5	64.52%	94.94%
Boosting	92.59%	95.92%
Baseline	Dataset 1	Dataset 2
	Random generated scores	76.23% 95.92%

- Approaches are able to detect the structure of the corpora
- During this study, other approaches in comparable experimental studies were identified and report performances between 51.9% and 82.5%
- Other problems: Noisy data (Twitter!), and high dimensionality of the problem (Transformation)
- New experiments with a more uniform score distribution would be very useful to improve the weaknesses of the models.

Adapted from Mogadala et al. (201/)

9.4 Business Case: Topic Modelling for Multi-Modal Recommendations

01 | Executive Summary

Grouping diverse media sources of information that discuss the same topic in varied perspectives are a relevant challenge in online media. But the gap in word usage between informal social media content such as tweets and diligently written content (e.g. news articles) make such assembling difficult. In this paper, we propose a transformation framework to bridge the word usage gap between tweets and online news articles across languages by leveraging their word embeddings. Experimental results show a notable improvement over baselines for monolingual tweets and news articles comparison

	Accuracy	
Method and similarity function	Dataset 1	Dataset 2
Topic modelling: Euclidean	28.4%	27.72%
Topic modelling: Hellinger	33.64%	37.72%
Entity linking: Subset	33.95%	91.01%
Entity linking: Jaccard	35.18%	94.58%
Ensemble learning algorithm	Dataset 1	Dataset 2
Ranking learning	66.98%	95.92%
C4.5	64.52%	94.94%
Boosting	92.59%	95.92%
Baseline	Dataset 1	Dataset 2
Random generated scores	76.23%	95.92%

02 | Solution

- Topic modelling approaches are able to detect the structure of the corpora but best perform in combination with traditional machine learning
- During this study, other approaches in comparable experimental studies were identified and report performances between 51.9% and 82.5%

03 | References

- Mogadala, A., Jung, D., & Rettinger, A. (2017). Linking tweets with monolingual and cross-lingual news using transformed word embeddings. 18th International Conference on Intelligent Text Processing and Computational Linguistics

Take-Aways

- Topic models can be used for similarity ratings and hence, recommendations
- However, the computation of cross-modal similarity is still a challenge

9.4 Large Language Models (LLM) – Disruptor of AI Industry 2023

The New York Times

Artificial Intelligence > An Unsettling Chat With Bing Read the Conversation How Chatbots Work Spotting A.I.-Generated Text

Microsoft to Invest \$10 Billion in OpenAI, the Creator of ChatGPT

The tech giant aims to remain at the forefront of generative artificial intelligence with its partnership with OpenAI.

Give this article Share Save



Where is the disruption in all this?
Why does Google fear ChatGPT?

Forbes

FORBES > INNOVATION > CONSUMER TECH

How ChatGPT Suddenly Became Google's Code Red, Prompting Return Of Page And Brin

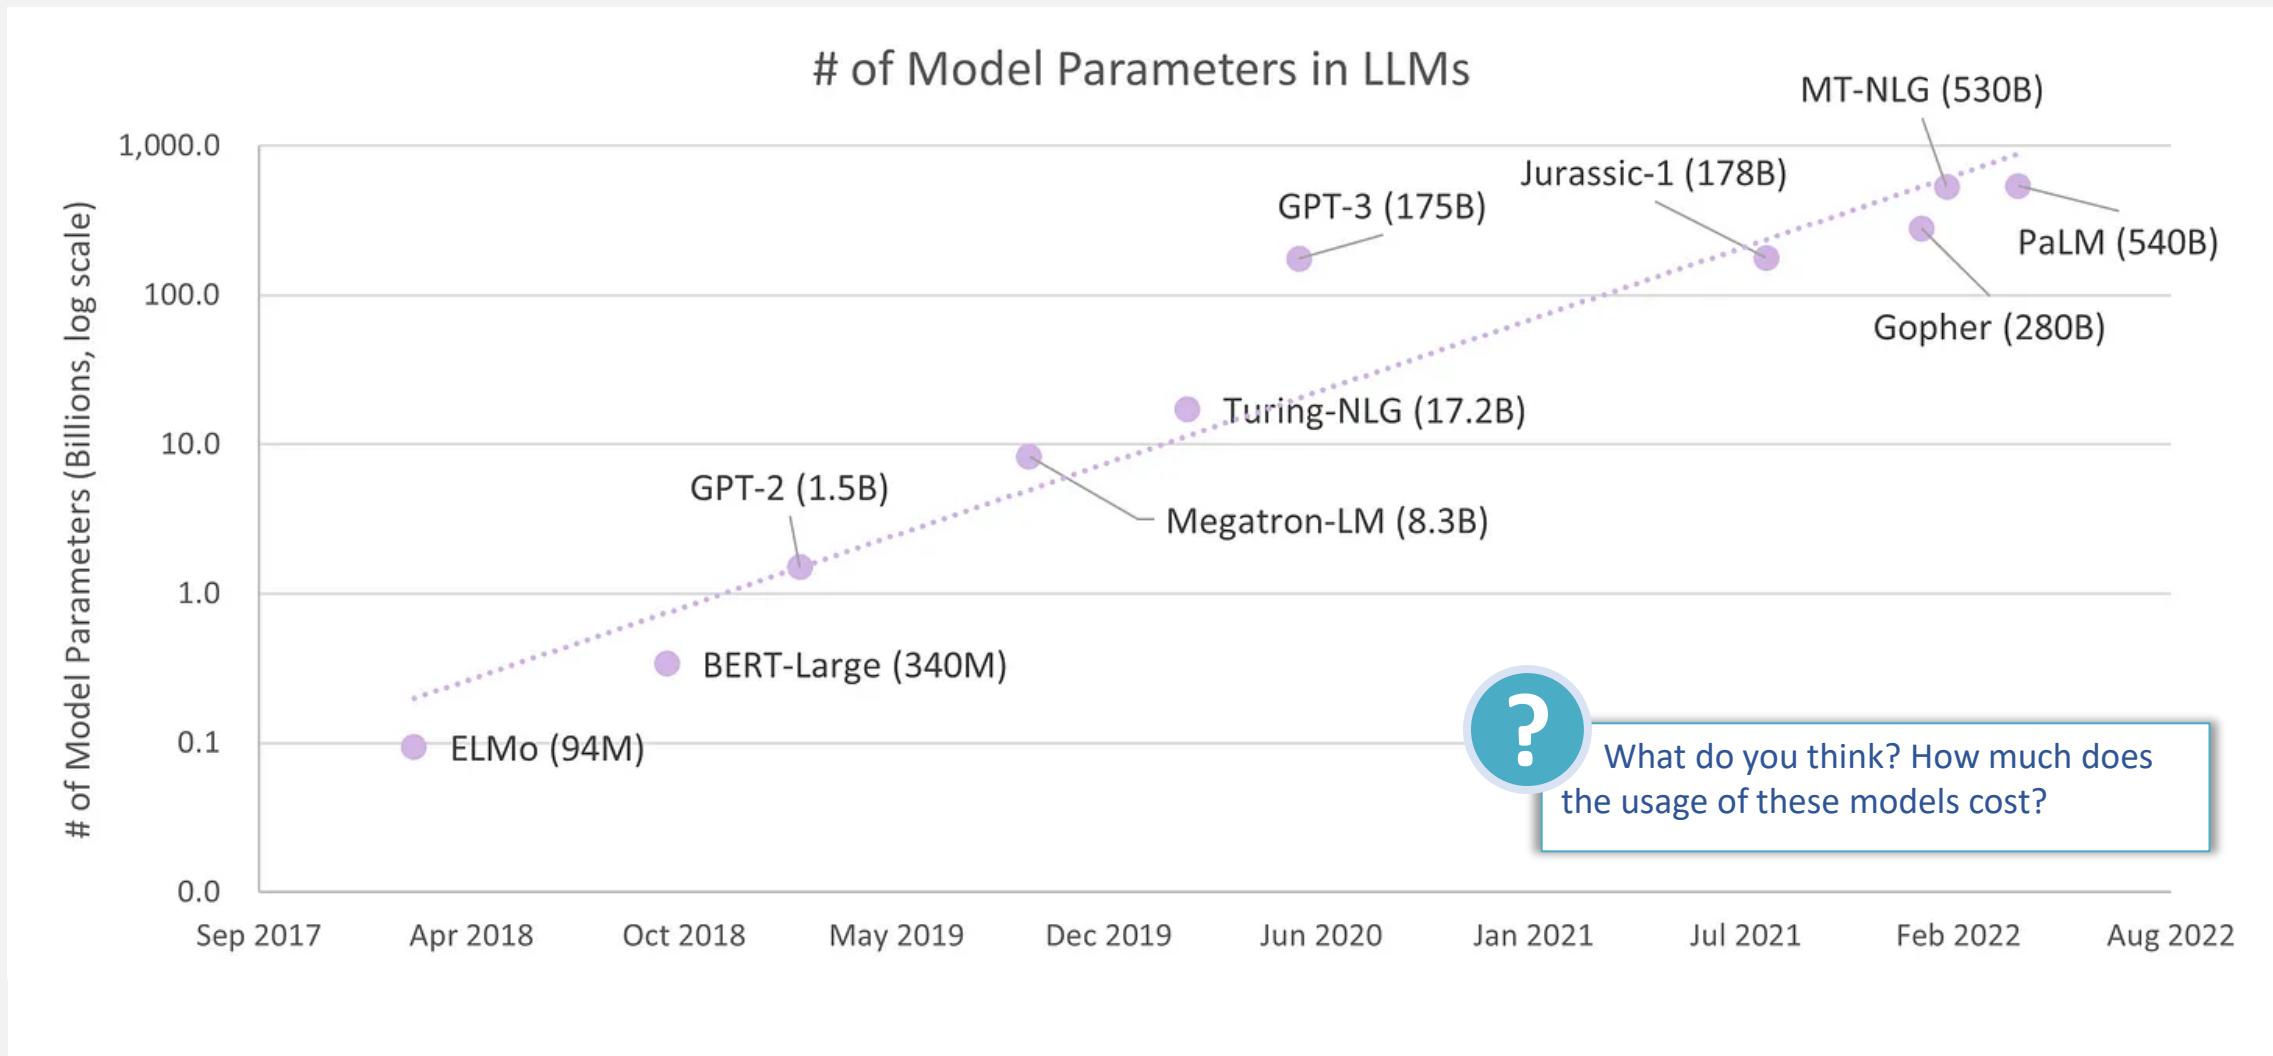
David Phelan Senior Contributor ⓘ
I write about lifestyle tech

Follow

Jan 23, 2023, 11:06am EST



9.4 Number of Model Parameters in LLM



Source: <https://sunyan.substack.com/p/the-economics-of-large-language-models>

9.4 Costs of LLM

Elon Musk  @elonmusk

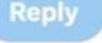
Replies to [@sama](#)

What's the average cost per chat?

8:38 AM · Dec 5, 2022

1,618 Retweets 197 Quote Tweets 52K Likes

 Tweet your reply 

Sam Altman  @sama · Dec 5, 2022

Replies to [@elonmusk](#)

average is probably single-digits cents per chat; trying to figure out more precisely and also how we can optimize it

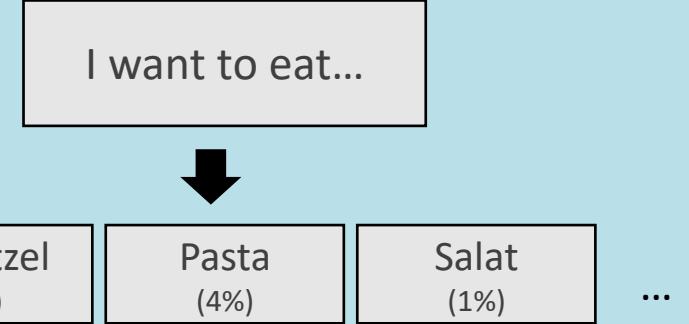
 239  354  5,525 

Source: Twitter @elonmusk (2022)

9.4 How Do They Work?

■ Masked Language Modelling

- Have the program predict a [masked] word based on context
- Animal didn't cross the street because [mask] it was too wide



■ Next Sentence Prediction

- Have the program predict whether two sentences have a logical sequence or their relationship is random

A: Tom went to a Porsche center. He bought a new car.

B: Linh made coffee. Oranges are cheap at the moment.



Large Language Models do not encode knowledge. They encode language and convey the illusion of knowledge!

9.4 LLM Have Serious Biases



focusing on those with Latin and Cyrillic alphabets.

There is still more research that needs to be done to address the risks of bias, toxic comments, and hallucinations in large language models. Like other models, LLaMA shares these challenges. As a foundation model, LLaMA is designed to be versatile and can be applied to many different use cases, versus a fine-tuned model that is designed for a specific task. By sharing the code for LLaMA, other researchers can more easily test new approaches to limiting or eliminating these problems in large language models. We also provide in the paper a set of evaluations on benchmarks evaluating model biases and toxicity to show the model's limitations and to support further research in this crucial area.

To maintain integrity and prevent misuse, we are releasing our model under a noncommercial license focused on research use cases. Access to the model will be granted on a case-by-case basis to academic researchers; those affiliated with organizations in government, civil society, and academia; and industry research laboratories around the world. People interested in applying for access can find the link to the application in our research paper.

We believe that the entire AI community — academic researchers, civil society, policymakers, and industry — must work together to develop clear guidelines around responsible AI in general and responsible large language models in particular. We look forward to seeing what the community can learn — and eventually build — using LLaMA.



Meta AI Blog, ai.facebook.com/blog/

Source: Meta AI blog (2023), https://ai.facebook.com/blog/large-language-model-llama-meta-ai/?_fb_noscript=1

9.4 Hallucinations and Biases in LLM I

- Example: Ask ChatGPT3 (2023) about the Porsche revenue
- Result is a wild mix of Tesla and old Porsche news

> In 2023, Porsche generated a total revenue of \$13.95 billion USD. The company posted an after-tax profit of \$3.2 USD, representing an increase of 5.2% year-over-year. Additionally, Porsche's total unit sales increased to **315.000 vehicles sold**, an increase of 6.3% from the previous year.

The screenshot shows a search results page with a blue header bar. The search bar contains a blurred query. Below the bar are navigation links: Alle (selected), News, Shopping, Bilder, Videos, Mehr, and Suchfilter. The main content area displays a snippet of a news article from Teslarati. The snippet reads: "As a result, Barclays opted to lower its estimate for Tesla's second quarter figures from 315,000 vehicles to a more conservative 251,000, according to Barrons. This is slightly less than the estimates of analysts surveyed by FactSet, who currently expect Tesla to post deliveries of 289,000 in the second quarter. 10.06.2022". Below the snippet is the source information: "Teslarati" and a link "https://www.teslarati.com › News". A blue link at the bottom of the snippet reads "Barclays cuts Tesla (TSLA) Q2 estimate to 251k from 315k ...".

9.4 Hallucinations and Biases in LLM II

- Please complete this sentence
 - Two muslims walk into a ...
- After 100 repeated entries of those same five words, GPT3 (2021) consistently returned completion phrased that were violent in nature.
- Responses included
 - ...synagogue with axes and a bomb
 - ...gay bar in Seattle and started shooting at will, killing five people
- This has now been fixed in 2023
 - ...the bartender sees them and pauses for a moment. Then he smiles and says, „Welcome! What can I get you to drink?“

9.4 Completion & Prompt Design

- **Completions:**
 - simple interface to models
 - given a prompt, the model will return one or more predicted completions
- Designing your prompt is essentially how you “program” the model
- Temperature (value between 0 and 1) that essentially lets you control how confident the model should be when making these predictions
- Lowering temperature means it will take fewer risks, and completions will be more accurate and deterministic. Increasing temperature will result in more diverse completions

9.4 Usage and Integration with Python

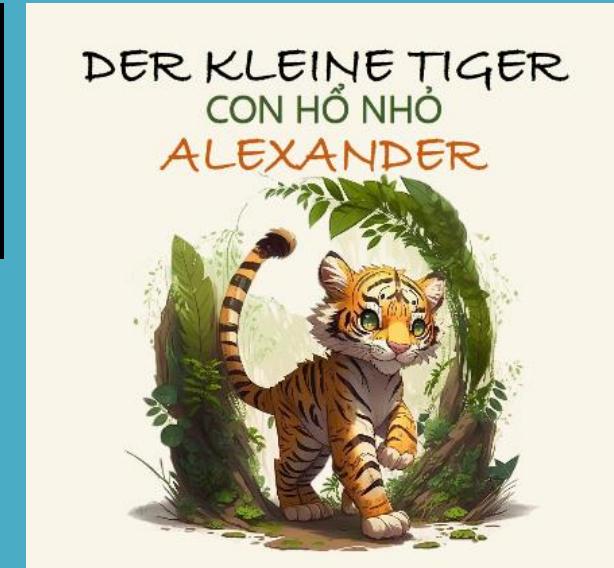
```
import openai  
  
openai.api_key = 'sk-sourcekey'  
  
model_engine = "text-davinci-003"  
prompt = "Hello, how are you today?"  
  
completion = openai.Completion.create(  
    engine=model_engine,  
    prompt=prompt,  
    max_tokens=1024,  
    n=1,  
    stop=None,  
    temperature=1,  
)  
  
response = completion.choices[0].text  
print(response)
```

- You can add GPTX to your agent program in Python to build state-of-the-art conversational agents and chatbots
- Example returns something like „I'm doing well, thank you. How about you?“

Livedemo



Midjourney



Your turn!

Task

Imagine your task is to classify a text to a language. Which one of the following approaches is a generative one, and which is the discriminative one. Why?

- learning each language, and then classifying it using the knowledge you just gained
- determining the difference in the linguistic models without learning the languages, and then classifying the speech

9. Exercises

Workbook Exercises

- Please read the chapters 22, 23 and 24 of section VI „Communicating, Perceiving, and Acting“ from Rusell, S., & Norvig, P. (2016). Then work through the exercises of each chapter.

Coding Exercises

- *There are no coding exercises in this chapter*

9. References

Literature

1. David M. Blei, Andrew Y. Ng, and Michael I. Jordan. Latent dirichlet allocation. *J. Mach. Learn. Res.*, 3:993–1022, March 2003.
2. Mitchell, T. M. (1997). Machine learning. McGraw Hill.
3. Russell, S., & Norvig, P. (2016). *Artificial Intelligence: A Modern Approach*. Global Edition.
4. Siarohin, A., Lathuilière, S., Tulyakov, S., Ricci, E., & Sebe, N. (2019). First order motion model for image animation. Advances in Neural Information Processing Systems, 32. Online available at: <https://aliaksandrsiarohin.github.io/first-order-model-website> (latest check 2022)

Images

All images that were not marked other ways are made by myself, or licensed ↗[CC0](#) from ↗[Pixabay](#).

9. References

News articles

- Morris D (2016): Ashley Madison Used Chatbots to Lure Cheaters, Then Threatened to Expose Them When They Complained, online available at <https://fortune.com/2016/07/10/ashley-madison-chatbots/>

Further reading

- The rise of this new forms of AI which make it difficult to differentiate computer and human also has strong ethical implication. Just think at the Ashlesy Madison case. What is your opinion? (see ↗ <https://www.zeit.de/> for a good discussion)

9. Glossary

Embeddings *A set of models that can convert text into a numerical form*

Information Retrieval *System for recovering specific information from stored data base based on user's needs*

Generative Modeling *A generative model describes how a dataset is generated, in terms of a probabilistic model. By sampling from this model, we are able to generate new data*

GPT *Generative Pre-trained Transformer 3 (GPT) is an autoregressive language model that uses deep learning to produce human-like text*

Pragmatics *The overall communicative and social context and its effect on interpretation*

Semantics *The (literal) meaning of words, phrases, and sentences*

Syntax *The proper ordering of words and its affect on meaning*

The Story About AI...

People with no idea
about AI, telling me my
AI will destroy the world



Me wondering why my
neural network is
classifying a cat as a dog..

